

Semiparametric SDF Estimators for Pooled, Non-Traded Cash Flows

Christian Tausch
AssetMetrix GmbH
Theresienhöhe 13, D-80339 Munich
christian.tausch@quant-unit.com

Alexander Bohnert
Munich University of Applied Sciences
Am Stadtpark 20, D-81243 Munich
alexander.bohnert@hm.edu

February 25, 2026

Keywords

Stochastic Discount Factor, Semiparametric Estimation, Spatial Dependence, Generalized Method of Moments (GMM), Private Equity, Illiquid Assets

Acknowledgements

We thank Nicolas Dütsch, Philipp Abel, and Kai Urban for valuable feedback and helpful comments that greatly improved the paper.

Declaration of interest

The authors report no conflict of interest. The authors alone are responsible for the content and writing of the paper.

Semiparametric SDF Estimators for Pooled, Non-Traded Cash Flows

Abstract

This paper analyzes stochastic discount factor estimation methodologies suited for pooled, non-traded cash flow streams such as the fund-level cash flows of private equity funds. The asymptotic inference framework for our semiparametric nonlinear least squares estimator draws on a spatial notion, i.e., the idea that the economic distance between distinct private equity funds can be measured. The empirical and Monte Carlo simulation results indicate (i) that our method can improve the popular Generalized Method of Moments approach of Driesssen et al. (2012), and (ii) that actively shifting the compounding horizon exploits a finite-sample bias–variance tradeoff stemming from a Timing Risk Premium. Our analysis further highlights the key trade-off between the time-series (number of vintage years) and cross-sectional (funds per vintage portfolio) dimensions, revealing that actively increasing the latter is highly effective in reducing estimator variance. Thus, we conjecture that traditional semiparametric extremum estimators like the one described by us shall be exclusively used for single-factor models until considerably more data for private equity funds becomes available.

1 Introduction

Private equity has outgrown its niche, sitting today on more than \$13 trillion in assets under management (McKinsey & Company, 2024), yet rigorous asset-pricing tools have not kept pace with this ascent. The empirical analysis and risk assessment of private equity and other non-traded cash flows remain fundamentally challenging due to the absence of market-based valuations and the inherent frictions of private markets (i.e., under incomplete information). Unlike public assets with trusted and tradeable valuations (in liquid secondary markets), private equity fund investments generate irregular, infrequently observed cash flows for which standard return-based asset pricing techniques are unsuitable.

To evaluate such irregular fund cash flows, the pioneering works of Driesssen et al. (2012) and Korteweg and Nagel (2016) introduced highly influential stochastic discount factor (SDF) estimators that have become the gold standard in private market asset pricing. The profound impact of their methodologies is evident from widespread adoption. The Driesssen et al. (2012) and Korteweg and Nagel (2016) estimators are heavily utilized to price a wide array of non-traded assets, spanning venture capital and real estate to infrastructure and collateralized loan obligations. Building on these major contributions, we propose a unified semiparametric framework that elegantly nests both estimators. The core benefit of our unified framework lies in providing the rigorous spatial asymptotic inference, missing in the original work of Driesssen et al. (2012), which ultimately allows researchers to identify, comprehend, and control the small-sample bias inherent to cash-flow-based SDF pricing.

Our nonlinear least squares estimator stems from the class of Least-Mean-Distance (LMD) estimators, which are easier to handle than traditional Generalized Method

of Moments (GMM) approaches (Pötscher and Prucha, 1997). Our LMD estimator arguably both simplifies and generalizes the GMM methodology of Driessen et al. (2012), where we provide the asymptotic inference framework that was missing in the original paper. The asymptotic inference formulations rely on the concept of spatial (near-epoch) dependence between funds as proposed by Korteweg and Nagel (2016). In this context, it is paramount to quantify the economic distance between funds by a measure like absolute vintage year difference or cash flow overlap¹.

For fund-level cash flow data of private equity funds, we document a small-sample bias term for cash-flow-based SDF estimators like Driessen et al. (2012) and Korteweg and Nagel (2016) that can be effectively decreased by increasing the cross-sectional sample size. The bias term arises due to a bias–variance tradeoff in finite samples. Using non-optimal discounting dates introduces an expected pricing error due to timing risk, but can substantially lower the estimator’s variance. Our estimator offers simple averaging over multiple discounting dates as one option to control (but not eliminate) the small-sample bias term.

In the empirical application of our new estimator, we evaluate simple linear and exponentially affine SDF models utilizing the well-known Preqin cash flow database across a grid of compounding horizons. While evaluating multi-factor models proves unreasonable given high estimator variance, our single-factor market models yield stable and economically meaningful estimates. At a bias-minimizing 15-year horizon, we estimate a market beta of 1.05 and 1.12 for equal- and value-weighted Private Equity portfolios, respectively. Crucially, while our rigorous asymptotic inference framework yields conservative standard errors, hv -block cross-validation confirms that these market beta estimates remain highly statistically significant. We document that the inherent estimator variance limiting multi-factor approaches can only be partially overcome by forming portfolios with a larger number of cross-sectional units per vintage year; highlighting a fundamental data constraint in modern private market asset pricing.

The paper is structured as follows. Section 1.1 reviews the related literature. Section 2 introduces our semiparametric LMD estimator and its corresponding asymptotic inference framework. Section 3 applies the method to estimate q^5 -investment-factor SDFs for various private equity fund types using simulated and real-world cash flows. Section 4 concludes. Appendix A analyzes the bias–variance tradeoff in the context of our SDF estimation problem. Appendix B compares our simulation and empirical results to their Driessen et al. (2012) counterparts.

1.1 Related Literature

For private equity funds, performance evaluation under the SDF framework remains an active research topic (Driessen et al., 2012; Korteweg and Nagel, 2016, 2024; Giommetti and Jørgensen, 2021; Gupta and Nieuwerburgh, 2021; Boyer et al., 2023; Korteweg, 2022) alongside an expanding list of applications across other private asset classes. Specifically, the Driessen et al. (2012) estimator is heavily utilized to price human capital contracts, venture capital, real estate and private equity funds (Kroencke et al., 2013; Kim et al., 2015; Farrelly and Stevenson, 2019; Lee et al.,

¹However, this economic inter-fund *distance* refers **not** to the term Least-Mean-*Distance* estimator.

2019; Zhao, 2020; Tausch and Pietz, 2024), while the Korteweg and Nagel (2016) approach has found broad application spanning private equity, infrastructure, real estate investment trusts, collateralized loan obligations, impact funds and private debt (Giommetti and Jørgensen, 2021; Andonov et al., 2021; Riddiough and Wiley, 2022; Boyer et al., 2023; Cordell et al., 2023; Erel et al., 2024; Jeffers et al., 2024; Korteweg and Nagel, 2024). Previously, SDF performance evaluation was already applied for public stock markets by, e.g., Farnsworth et al. (2002) and for hedge fund returns by Li et al. (2016).

Within the private equity literature, approaches can be broadly classified into parametric and semiparametric methods. Parametric models, such as those proposed by Franzoni et al. (2012) and Ang et al. (2018), explicitly specify the return dynamics and the underlying data generating process of private equity investments. While these methods can uncover latent factors or deal-level dynamics, they typically rely on strong distributional assumptions (e.g., log-normality) that may not fully capture the irregular and skewed nature of PE cash flows.

Consequently, semiparametric approaches have gained prominence, starting with the seminal cross-sectional GMM framework of Driessen et al. (2012). They estimate a linear factor model that minimizes the squared pricing errors of PE vintage year portfolios, explicitly evaluating cash flows at fund inception. Building on this, Korteweg and Nagel (2016) introduced the Generalized Public Market Equivalent (GPME) framework, shifting the estimation of the SDF to public market data while evaluating PE cash flows in a second step. Crucially, Korteweg and Nagel (2016) were the first to apply a spatial heteroskedasticity and autocorrelation consistent (SHAC) covariance matrix estimator (Conley, 1999) to account for the economic distance (cash flow overlap) between PE funds, moving beyond simple cross-sectional bootstrapping. Further complementing this cash flow dimension, a strand of literature models the informational frictions and smoothing inherent in reported Net Asset Values (NAVs) (Brown et al., 2019; Jenkinson et al., 2020; Brown et al., 2021). Moreover, Hüther et al. (2022) leverage a similar generalized framework to evaluate private firms using credit market equivalents.

Our paper’s unified LMD framework contributes to this semiparametric literature by formally nesting both the Driessen et al. (2012) and Korteweg and Nagel (2016) approaches while supplying the rigorous increasing-domain asymptotic inference missing in earlier work (drawing on the spatial theory of Pötscher and Prucha (1997), Kim and Sun (2011), and Jenish and Prucha (2012)). Furthermore, while Gredil et al. (2020) highlighted that NPV-based measures are biased relative to per-period returns, we explicitly formalize this discrepancy as a “Timing Risk Premium”—a concept conceptually linked to recent findings on discount-rate risk in secondary markets (Boyer et al., 2023).

2 Methodology

We now formalize the semiparametric estimator that synthesizes the approaches of Driessen et al. (2012) and Korteweg and Nagel (2016). A detailed exposition of how our unified framework departs from and improves upon these base models is provided in Section 2.5.

In a nutshell, the Driessen et al. (2012) estimator determines the estimated alpha and beta parameters of a linear factor model by minimizing the squared net present value (NPV) of all n fund cash flows in a least-squares optimization problem.

$$(\hat{\alpha}, \hat{\beta}) = \arg \min_{\alpha, \beta} \sum_{i=1}^n [NPV_i(\alpha, \beta)]^2 \quad (1)$$

with NPV for the i th fund

$$NPV_i(\alpha, \beta) = \sum_{t \in \{t_{0,i}, t_{0,i}+1, \dots\}} \frac{\text{CashFlow}_{i,t}}{\prod_{s=t_{0,i}+1}^t (1 + r_{f,s} + \alpha + \beta r_{\text{market},s})}$$

Here, the Driessen et al. (2012) approach relies on the identifying assumption that the model is correctly specified, so that $\mathbb{E}[NPV_i(\theta_0)] = 0$.

In this Methodology section, we will show:

1. Why this estimator yields conditionally expected pricing errors of zero under a Net Present Value framework, and we analytically characterize the bias term (timing risk premium) introduced by a Net Future Value approach (Section 2.1).
2. How to modify the estimator to be able to control the bias term (Section 2.2).
3. How we can estimate asymptotic standard errors for the parameters alpha and beta (Sections 2.3 and 2.4).
4. How our estimator compares to similar approaches (Section 2.5.)

2.1 Asset Pricing for Pooled Cash Flows

Let private equity fund $i = 1, 2, \dots, n$ be characterized by its net cash flows $CF_{i,t}$ (i.e., distributions minus contributions) and its net asset values $NAV_{i,t}$ with discrete time index $t = 0, 1, 2, \dots, T$. To increase the mathematical tractability of the problem, we assume a deal-by-deal data generating process (DGP) for CF where each fund deal consists exactly of one investment and one divestment cash flow in combination with a simple return model for the multi-period deal returns. This means the fund-level cash flow process $(CF_{i,t})_{t=0,1,\dots,T}$ is an aggregation of deal-level cash flow pairs consisting of one negative at deal inception and at least one positive cash flow later $CF_{i,t} = \sum_j^J cf_{j,i,t}$.

Assumption 1. *Deal-level data generation process:*

1. *Each fund i consists of J underlying deals.*
2. *Each deal is characterized by exactly one, negative investment cash flow, denoted by $\text{Inv}_{i,j}$, which occurs at time $t_{i,j}^{\text{Inv}} \in \{0, 1, \dots, T - 1\}$. It holds $\text{Inv}_{i,j} < 0$.*
3. *Each deal is characterized by a positive divestment cash flow stream, denoted by $(\text{Div}_{i,j,k})_{k=1,\dots,K_{i,j}}$, which occur after the investment cash flow $t_{i,j,k}^{\text{Div}} > t_{i,j}^{\text{Inv}}$ for all k . It holds $\text{Div}_{i,j,k} > 0$.*
4. *The cumulative fund cash flows are generated by summarizing over all deal-level cash flows, i.e., $\sum_{t=0}^T CF_{i,t} = \sum_{j=1}^J (\text{Inv}_{i,j} + \sum_{k=1}^{K_{i,j}} \text{Div}_{i,j,k})$ for all i .*

Symbol	Description
n	Number of cross-sectional units (funds or portfolios)
V	Number of distinct vintage years
n_v	Number of portfolios in vintage year v
v_i	Vintage year of fund/portfolio i
J	Number of deals per fund
$K_{i,j}$	Number of divestment cash flows for deal j of fund i
$CF_{i,t}$	Net cash flow of fund i at time t
Ψ_t, M_t	Multi-period and single-period SDF
$\Psi_{\tau,t}(\theta)$	SDF ratio Ψ_t/Ψ_τ , parameterized by θ
$\theta = (\alpha, \beta)$	SDF parameter vector (intercept and factor loadings)
θ_0	True parameter vector
$\delta_{i,j}$	Deal-level pricing error
$\epsilon_{\tau,i}$	Fund pricing error at discounting date τ
$\bar{\epsilon}_i$	Weighted average pricing error of fund i
a_i	Fund-level alpha (abnormal return; distinct from SDF α)
\mathcal{T}_i	Set of discounting dates for fund i
w_i	Weighting factor for fund i
$L(\cdot)$	Loss function (e.g., $L(x) = x^2$)
$Q_n(\theta)$	Sample objective function
$d_{i,j}$	Economic distance between units i and j
$D (= b_n)$	SHAC bandwidth parameter
H	Hessian of population objective at θ_0
Λ	Long-run covariance matrix of the score
Σ	Asymptotic covariance matrix of $\hat{\theta}$

Table 1: Key notation for the Methodology section.

From asset pricing theory, we know that we can use a stochastic discount factor Ψ_t to price each underlying deal

$$\mathbb{E} \left[\delta_{i,j} \mid \mathcal{F}_{t_{i,j}^{\text{Inv}}} \right] = 0 \quad \forall i, j \quad (2)$$

where we denote the deal-level pricing error by

$$\delta_{i,j} := \text{Inv}_{i,j} + \sum_{k=1}^{K_{i,j}} \frac{\Psi_{t_{i,j,k}^{\text{Div}}}}{\Psi_{t_{i,j}^{\text{Inv}}}} \text{Div}_{i,j,k} \quad (3)$$

Before we can derive the expected value for the fund-level pricing errors, let us analyze the deal-level counterparts.

Case A: Discounting deals to a date before investment ($\tau < t_{i,j}^{\text{Inv}}$)
First we apply the Law of Iterated Expectations (the so-called “tower property”) to determine the fair value of individual deal cash flows at time $\tau < t_{i,j}^{\text{Inv}}$.

$$\mathbb{E} \left[\frac{\Psi_{t_{i,j}^{\text{Inv}}}}{\Psi_\tau} \delta_{i,j} \mid \mathcal{F}_\tau \right] = \mathbb{E} \left[\mathbb{E} \left[\frac{\Psi_{t_{i,j}^{\text{Inv}}}}{\Psi_\tau} \delta_{i,j} \mid \mathcal{F}_{t_{i,j}^{\text{Inv}}} \right] \mid \mathcal{F}_\tau \right] \quad (4)$$

Since $\tau < t_{i,j}^{\text{Inv}}$, it holds that $\mathcal{F}_\tau \subseteq \mathcal{F}_{t_{i,j}^{\text{Inv}}}$. Furthermore, the term $\frac{\Psi_{t_{i,j}^{\text{Inv}}}}{\Psi_\tau}$ is fully known at time $t_{i,j}^{\text{Inv}}$, so it can be pulled out of this inner expectation to receive

$$\mathbb{E} \left[\frac{\Psi_{t_{i,j}^{\text{Inv}}}}{\Psi_\tau} \delta_{i,j} \mid \mathcal{F}_\tau \right] = \mathbb{E} \left[\underbrace{\frac{\Psi_{t_{i,j}^{\text{Inv}}}}{\Psi_\tau} \mathbb{E} [\delta_{i,j} \mid \mathcal{F}_{t_{i,j}^{\text{Inv}}}] \mid \mathcal{F}_\tau}_{=0 \text{ by Eq. 2}} \right] = \mathbb{E} \left[\frac{\Psi_{t_{i,j}^{\text{Inv}}}}{\Psi_\tau} \cdot 0 \mid \mathcal{F}_\tau \right] = 0 \quad (5)$$

Case B: Discounting deals to a date after investment ($\tau > t_{i,j}^{\text{Inv}}$) When the discounting/compounding date τ is after the investment date, the term $\Psi_{t_{i,j}^{\text{Inv}}} / \Psi_\tau$ is not known at $t_{i,j}^{\text{Inv}}$ (it depends on future SDF realizations). Thus, the previous logic does not hold. Instead, we have a non-zero term that represents the accumulated realized pricing error.

Let $\tau > t_{i,j}^{\text{Inv}}$. The expectation becomes:

$$\mathbb{E} \left[\frac{\Psi_{t_{i,j}^{\text{Inv}}}}{\Psi_\tau} \delta_{i,j} \mid \mathcal{F}_\tau \right] = \frac{\Psi_{t_{i,j}^{\text{Inv}}}}{\Psi_\tau} \mathbb{E} [\delta_{i,j} \mid \mathcal{F}_\tau] \quad (6)$$

Since we condition on \mathcal{F}_τ and $t_{i,j}^{\text{Inv}} < \tau$, the SDF ratio is known and can be pulled out. Then the term $\mathbb{E} [\delta_{i,j} \mid \mathcal{F}_\tau]$ is the conditional expectation of the pricing error given information available at time τ . It decomposes as:

$$\mathbb{E} [\delta_{i,j} \mid \mathcal{F}_\tau] = \underbrace{\text{Inv}_{i,j} + \sum_{t_{i,j,k}^{\text{Div}} \leq \tau} \frac{\Psi_{t_{i,j,k}^{\text{Div}}}}{\Psi_{t_{i,j}^{\text{Inv}}}} \text{Div}_{i,j,k}}_{\text{Realized Cash Flows}} + \underbrace{\mathbb{E} \left[\sum_{t_{i,j,k}^{\text{Div}} > \tau} \frac{\Psi_{t_{i,j,k}^{\text{Div}}}}{\Psi_{t_{i,j}^{\text{Inv}}}} \text{Div}_{i,j,k} \mid \mathcal{F}_\tau \right]}_{\text{Expected Future Cash Flows}} \quad (7)$$

Generally, this sum can become (is usually) non-zero because it includes realized shocks (abnormal returns) that occurred between $t_{i,j}^{\text{Inv}}$ and τ .

Case C: Compounding deals to a future date Let us analyze the following Net Future Value (NFV) case, where $t_{i,j}^{\text{Future}}$ denotes a compounding target date after the deal investment date (i.e., $\tau < t_{i,j}^{\text{Inv}} < t_{i,j}^{\text{Future}}$).

$$\begin{aligned} \mathbb{E} \left[\frac{\Psi_{t_{i,j}^{\text{Future}}}}{\Psi_\tau} \delta_{i,j} \mid \mathcal{F}_\tau \right] &= \mathbb{E} \left[\frac{\Psi_{t_{i,j}^{\text{Inv}}}}{\Psi_\tau} \frac{\Psi_{t_{i,j}^{\text{Future}}}}{\Psi_{t_{i,j}^{\text{Inv}}}} \delta_{i,j} \mid \mathcal{F}_\tau \right] \\ &= \mathbb{E} \left[\frac{\Psi_{t_{i,j}^{\text{Inv}}}}{\Psi_\tau} \mathbb{E} \left[\frac{\Psi_{t_{i,j}^{\text{Future}}}}{\Psi_{t_{i,j}^{\text{Inv}}}} \delta_{i,j} \mid \mathcal{F}_{t_{i,j}^{\text{Inv}}} \right] \mid \mathcal{F}_\tau \right] \\ &= \mathbb{E} \left[\frac{\Psi_{t_{i,j}^{\text{Inv}}}}{\Psi_\tau} \text{Cov} \left(\frac{\Psi_{t_{i,j}^{\text{Future}}}}{\Psi_{t_{i,j}^{\text{Inv}}}}, \delta_{i,j} \mid \mathcal{F}_{t_{i,j}^{\text{Inv}}} \right) \mid \mathcal{F}_\tau \right] \end{aligned} \quad (8)$$

The term can generally become non-zero, representing the covariance between the deal's pricing error and the future level of the SDF, effectively a risk premium for the timing mismatch. Note that the final line of Equation 8 is leveraging

$$\mathbb{E} \left[\frac{\Psi_{t_{i,j}^{\text{Future}}}}{\Psi_{t_{i,j}^{\text{Inv}}}} \delta_{i,j} \mid \mathcal{F}_{t_{i,j}^{\text{Inv}}} \right] = \text{Cov} \left(\frac{\Psi_{t_{i,j}^{\text{Future}}}}{\Psi_{t_{i,j}^{\text{Inv}}}}, \delta_{i,j} \mid \mathcal{F}_{t_{i,j}^{\text{Inv}}} \right) + \mathbb{E} \left[\frac{\Psi_{t_{i,j}^{\text{Future}}}}{\Psi_{t_{i,j}^{\text{Inv}}}} \mid \mathcal{F}_{t_{i,j}^{\text{Inv}}} \right] \underbrace{\mathbb{E} \left[\delta_{i,j} \mid \mathcal{F}_{t_{i,j}^{\text{Inv}}} \right]}_{=0 \text{ by Eq. 2}} \quad (9)$$

This result highlights a fundamental dichotomy. For *ex-ante* asset pricing tests, the Net Future Value approach is misspecified: the non-zero conditional covariance implies that the conditional moment condition $\mathbb{E}[m_{i,j}^{\text{NFV}}(\theta_0) \mid \mathcal{F}_\tau] = 0$ is violated, leading to biased parameter estimates if not explicitly corrected for this risk premium. Conversely, for *ex-post* performance benchmarking, this compounding approach remains valuable. By compounding realized cash flows to a future date T , we facilitate an intuitive comparison of the fund's generated terminal wealth against a public benchmark, provided one acknowledges that the resulting *alpha* implicitly includes the compensation for the timing risk identified above. This entails that any *alpha* measured via Net Future Values is mechanically a composite metric. Specifically, we can mathematically decompose the expected NFV pricing error into two distinct components:

$$\underbrace{\mathbb{E} \left[\frac{\Psi_{t_{i,j}^{\text{Future}}}}{\Psi_\tau} \delta_{i,j} \mid \mathcal{F}_\tau \right]}_{\text{Measured Alpha}} = \underbrace{\mathbb{E} \left[\frac{\Psi_{t_{i,j}^{\text{Inv}}}}{\Psi_\tau} \frac{a_{i,j}^{\text{True}}}{R_{f,t_{i,j}^{\text{Inv}} \rightarrow t_{i,j}^{\text{Future}}}} \mid \mathcal{F}_\tau \right]}_{\text{True Skill}} + \underbrace{\mathbb{E} \left[\frac{\Psi_{t_{i,j}^{\text{Inv}}}}{\Psi_\tau} \text{Cov} \left(\frac{\Psi_{t_{i,j}^{\text{Future}}}}{\Psi_{t_{i,j}^{\text{Inv}}}}, \delta_{i,j} \mid \mathcal{F}_{t_{i,j}^{\text{Inv}}} \right) \mid \mathcal{F}_\tau \right]}_{\text{Timing Risk Premium}} \quad (10)$$

where R_f is the gross risk-free return over the holding period and a^{True} is the deal's expected pricing error at inception (fund-level abnormal return, distinct from the SDF intercept α). Even if a manager has no skill ($a^{\text{True}} = 0$), the measured alpha will include the *Timing Risk Premium* which is the (discounted) covariance between the future SDF realization and the deal valuation. This is not a flaw but a feature of *ex-post* benchmarking: it correctly reflects that the investor was exposed to this specific path of SDF realizations. However, for *ex-ante* pricing tests, the presence of this non-zero *Timing Risk Premium* is problematic. Standard econometric estimators are designed to find parameters that set the average pricing error to zero. If the theoretical moment contains this positive premium, the estimator will be forced to compensate by biasing the estimated alpha downwards or distorting the SDF parameters to "explain away" the *Timing Risk Premium*.

Fund-level pricing error For the pooled, fund-level cash flow stream, we assume that the true fund valuation $V_{i,\tau}$ is generally not observable for us

$$V_{i,\tau} := \mathbb{E} \left[\sum_{t=\tau}^T \frac{\Psi_t}{\Psi_\tau} CF_{i,t} \mid \mathcal{F}_\tau \right] \quad \forall \tau \geq \min_j t_{i,j}^{\text{Inv}} \quad (11)$$

For active funds with $\tau > \min_j t_{i,j}^{\text{Inv}}$, we merely can observe reported fund NAVs as proxies for the true fund value, which are known to suffer from stale pricing and appraisal-smoothing problems (Brown et al., 2019; Korteweg, 2019; Jenkinson et al., 2020)². Additionally, we define the vintage year for each fund in a mathematically convenient way as its “starting year” $v_i := \text{YEAR}(\min_j t_{i,j}^{\text{Inv}})$.

For any discounting date $\tau \leq \min_j t_{i,j}^{\text{Inv}}$, the expected fund pricing error is equal to zero.

Proposition 1. (*Price of a pooled cash flow stream at fund inception*)

$$\mathbb{E} \left[\sum_{t=\tau}^T \frac{\Psi_t}{\Psi_\tau} CF_{i,t} | \mathcal{F}_\tau \right] = 0 \quad (12)$$

if $\tau \leq \min_j t_{i,j}^{\text{Inv}}$.

Proof. Let us first analyze the case $\tau \leq \min_j t_{i,j}^{\text{Inv}}$ by expressing the fund cash flows in term of discounted deal pricing errors using Equation 3.

$$\mathbb{E} \left[\sum_{t=\tau}^T \frac{\Psi_t}{\Psi_\tau} CF_{i,t} | \mathcal{F}_\tau \right] = \mathbb{E} \left[\sum_{j=1}^J \frac{\Psi_{t_{i,j}^{\text{Inv}}}}{\Psi_\tau} \delta_{i,j} | \mathcal{F}_\tau \right] \quad (13)$$

Then we can simply plug in Equation 5 to receive the postulated zero expected fund pricing error.

$$\mathbb{E} \left[\sum_{j=1}^J \frac{\Psi_{t_{i,j}^{\text{Inv}}}}{\Psi_\tau} \delta_{i,j} | \mathcal{F}_\tau \right] = \mathbb{E} \left[\sum_{j=1}^J \frac{\Psi_{t_{i,j}^{\text{Inv}}}}{\Psi_\tau} \cdot 0 | \mathcal{F}_\tau \right] = 0 \quad (14)$$

□

Next, let us analyze the case $\tau > \min_j t_{i,j}^{\text{Inv}}$ for Equation 11. For any deal j with $t_{i,j}^{\text{Inv}} < \tau$, the conditional expectation $\mathbb{E}[\delta_{i,j} | \mathcal{F}_\tau]$ represents the realized abnormal return of the deal up to time τ . For a generic risky asset, this realized return is a random variable with non-zero variance. Therefore, the condition that Equation 11 equals 0 holds *ex-post* only by chance (i.e., if realized shocks cancel out exactly). Ex-ante, without knowledge of realizations, $\tau \leq \min_j t_{i,j}^{\text{Inv}}$ is the necessary condition to guarantee a zero expected pricing error.

2.2 Least-Mean-Distance estimator

In this subsection, we introduce a new SDF estimator designed to analyze the effect of different discounting dates τ on the bias and variance of the estimated SDF parameters. In the previous subsection, we demonstrated that only the fund inception date yields unconditionally unbiased parameter estimates in a NPV-style optimization formulation. However, to be able to empirically test the impact of

²In the empirical section, we will treat the most recent NAV as the final distribution cash flow for non-liquidated funds to increase the sample data.

mixing in biased NFV terms, we choose to generalize our estimation framework so we can handle both cases flexibly. With this respect, our general idea is to rather average over multiple “suitable” discounting date candidates τ than to select only one candidate for the “unconditionally optimal” discounting date.

Henceforth, we assume that the underlying transactions within a private equity fund cannot be distinguished individually, and that only the funds total (pooled) cash flows are observable. The stochastic discount factor $\Psi_{\tau,t}$ is used to calculate the time- τ “price” $P_{\tau,t,i}$ of a **single** time- t cash flow of any given PE fund i

$$P_{\tau,t,i} := \Psi_{\tau,t} \cdot CF_{i,t} = \frac{\Psi_t}{\Psi_\tau} \cdot CF_{i,t} \quad \forall \tau, t, i \quad (15)$$

with multi-period SDF $\Psi_t = \prod_{k=1}^t M_k$ and $M_0 = 1$. As SDFs are commonly parameterized by a vector $\theta \in \mathbb{R}^p$, i.e., $\Psi_{\tau,t} \equiv \Psi_{\tau,t}(\theta)$, our goal is to find an estimation method for the optimal θ . We denote this optimal/best/true parameter vector as θ_0 . We call the numerator Ψ_t the discount part of the multi-period SDF $\Psi_{\tau,t}$ (used for present value calculations) and the denominator Ψ_τ the compound part (used for future value calculations). For each fund i and all points τ within a common fund lifetime, the empirical “pricing error” $\epsilon_{\tau,i}$ of **all** fund cash flows is calculated as time- τ “present value”

$$\epsilon_{\tau,i} := \sum_{t=0}^T P_{\tau,t,i} \quad \forall \tau, i \quad (16)$$

We use the terms “price”, “pricing error” and “net present value” in quotation marks to acknowledge the theoretical asset pricing problem which can arise for pooled cash flows and has been described in the previous subsection.

To better analyze the impact of different discount date τ on the estimator’s bias and variance, we define the (w_i -weighted) average pricing error $\bar{\epsilon}_i$ that averages over the set \mathcal{T}_i

$$\bar{\epsilon}_i = w_i \cdot \frac{1}{\text{card}(\mathcal{T}_i)} \sum_{\tau \in \mathcal{T}_i} \epsilon_{\tau,i} \quad \forall i \quad (17)$$

where \mathcal{T}_i gives the set of discounting dates τ for fund i which is more thoroughly described below. Finally, the scalar weighting factor w_i can be (i) one divided by the fund’s invested capital for equal weighting of funds, (ii) one divided by the vintage year sum of invested capital for vintage year weighting, (iii) the scalar one for fund-size weighting, or (iv) some macroeconomic deflator.

The role of \mathcal{T}_i The set \mathcal{T}_i controls which discounting dates enter the averaged pricing error $\bar{\epsilon}_i$ in Equation 17, as visualized in Figure 1. The smallest choice $\mathcal{T}_i = \{\min_j t_{i,j}^{\text{Inv}}\}$ corresponds to a classical NPV calculation and is used by both Driessens et al. (2012) and Korteweg and Nagel (2016). The largest candidate set includes all periods from inception to the present.

Averaging over multiple dates in \mathcal{T}_i was originally motivated by the “exploding alpha” problem first described by Driessens et al. (2012): when the initial cash flow is very small, an arbitrarily large alpha can discount all subsequent cash flows close

to zero, rendering the beta factors irrelevant and distorting estimation. Averaging over \mathcal{T}_i mitigates this degeneracy.

However, as shown in Section 2.1, including future value dates in \mathcal{T}_i introduces a timing risk premium bias. The bias–variance tradeoff of different $|\mathcal{T}_i|$ choices is studied empirically in the simulation study (Section 3.3).

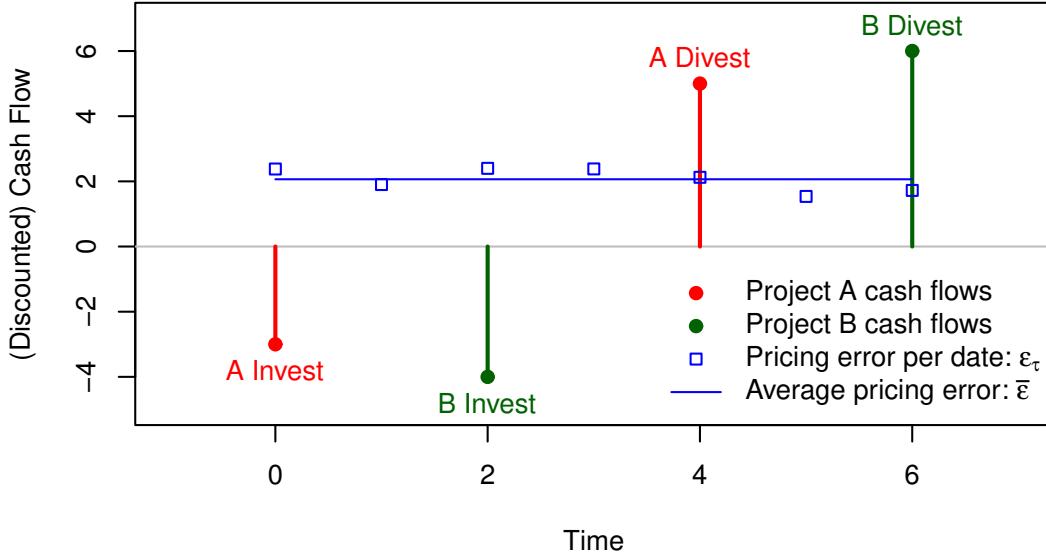


Figure 1: How to calculate and interpret the average pricing error introduced by Equation 17: The time index t is relevant for the net cash flows (black dots). The time index τ is used for the pricing error, i.e., the sum of discounted net cash flows (blue boxes). The weighted average of these "pricing errors" gives the average pricing error $\bar{\varepsilon}$ as defined in Equation 17 (solid blue line). In this example, $\text{card}(\mathcal{T}_i) = 7$, i.e., the number of blue boxes.

Assumption 2. *To guarantee an unconditionally unbiased estimator, we assume to always only discount to the fund start date $\mathcal{T}_i = \left\{ \min_j t_{i,j}^{\text{Inv}} \right\}$.*

Remark 1. Assumption 2 guarantees that $\mathbb{E}[\bar{\varepsilon}_i(\theta_0)] = 0$ under the true parameter. For the quadratic loss $L(x) = x^2$, the population objective $Q_0(\theta) = -\mathbb{E}[\bar{\varepsilon}_i(\theta)^2] = -\text{Var}[\bar{\varepsilon}_i(\theta)] - (\mathbb{E}[\bar{\varepsilon}_i(\theta)])^2$ involves both variance and bias terms. Consistency then requires that θ_0 uniquely minimizes this composite criterion (see Proposition 2).

To find the optimal value for θ , we select an estimator from the broad class of extremum estimators.

Definition 1. Extremum estimator (Newey and McFadden, 1994, Equation 1.1): An estimator $\hat{\theta}$ is an extremum estimator if there is an objective function $Q_n(\theta)$ such that

$$\hat{\theta} = \arg \max_{\theta} Q_n(\theta)$$

for $\theta \in \Theta$ where Θ is the set of all possible parameter values.

Concretely, our Least Mean Distance (LMD) estimator (Pötscher and Prucha, 1997, Equation 7.1) minimizes the average loss of $\bar{\epsilon}$. For consistency with Definition 1, we will technically maximize the average negative loss of $\bar{\epsilon}$, which is equivalent.

$$\hat{\theta} = \arg \min_{\theta \in \Theta} \frac{1}{n} \sum_{i=1}^n L(\bar{\epsilon}_i) \equiv \arg \max_{\theta \in \Theta} Q_n(\theta) \quad \text{with} \quad Q_n(\theta) = -\frac{1}{n} \sum_{i=1}^n L(\bar{\epsilon}_i) \quad (18)$$

where $L : \mathbb{R}^1 \rightarrow \mathbb{R}^1$ denotes a loss function, e.g., $L(x) = (x - 0)^2$ in the case of nonlinear least squares. Throughout the paper, the weighted average fund pricing error $\bar{\epsilon}_i \equiv \bar{\epsilon}_i(\theta)$ is regarded as nonlinear random function of the SDF parameter θ .

2.3 Asymptotic framework

By the value-additivity principle (Hansen and Richard, 1987), our SDF model holds for any aggregation level of cash flows. In practice, we form vintage-year portfolios as cross-sectional units, which mitigates idiosyncratic noise and GP financial engineering artifacts while preserving sufficient test assets for parameter identification.

Assumption 3. *For each vintage year, we pool fund cash flows to form n_v portfolios. Thus, $n = \sum_{v=1}^V n_v$. The boundary cases are (i) single-fund portfolios and (ii) one portfolio per vintage year. Both equal-weighted and fund-value-weighted portfolios are considered.*

Without loss of generality, we refer to cross-sectional units as funds. The detailed comparison of aggregation levels is provided in Section 3.3.

The pricing error $\bar{\epsilon}_i$ is naturally indexed by both fund identity and vintage year, yielding a random field $\bar{\epsilon}_i \equiv \bar{\epsilon}_{i,v_i}$ that fits neither standard time-series, cross-sectional, nor panel data frameworks (cf. Figure 2). We address this by adopting a spatial framework similar to Korteweg and Nagel (2016), where they measure the economic distance between funds as a function of the cash flow overlap. Thus, we generally consider $\bar{\epsilon}$ from Equation 17 as a random field (cf. Figure 2). In our case, it is convenient to interpret the fund vintage year v_i as second dimension in our pricing error random field, i.e., $\bar{\epsilon}_i \equiv \bar{\epsilon}_{i,v_i}$.

Yet, in this section, we mainly follow the time-series asymptotic framework of Pötscher and Prucha (1997) since our “spatial” distance measure (between vintage years) is time, and adaptation to our case is thus straightforward. If we observe only one fund per vintage year (or, equivalently, form vintage year portfolios), we will easily see that the framework of Pötscher and Prucha (1997) with time-series indexing can be applied without any major modification.

2.3.1 Vintage year asymptotics

We assume that the “spatial” (i.e., economic) distance between cross-sectional units, i.e., private equity funds/portfolios, can be measured quantitatively³. Our “cross-sectional” asymptotic theory lets the number of funds go to infinity $n \rightarrow \infty$. To

³Generally, the economic distance measure could include multiple dimensions, e.g., temporal, geographic, and industry sector proximity. This could be an interesting topic for future research.

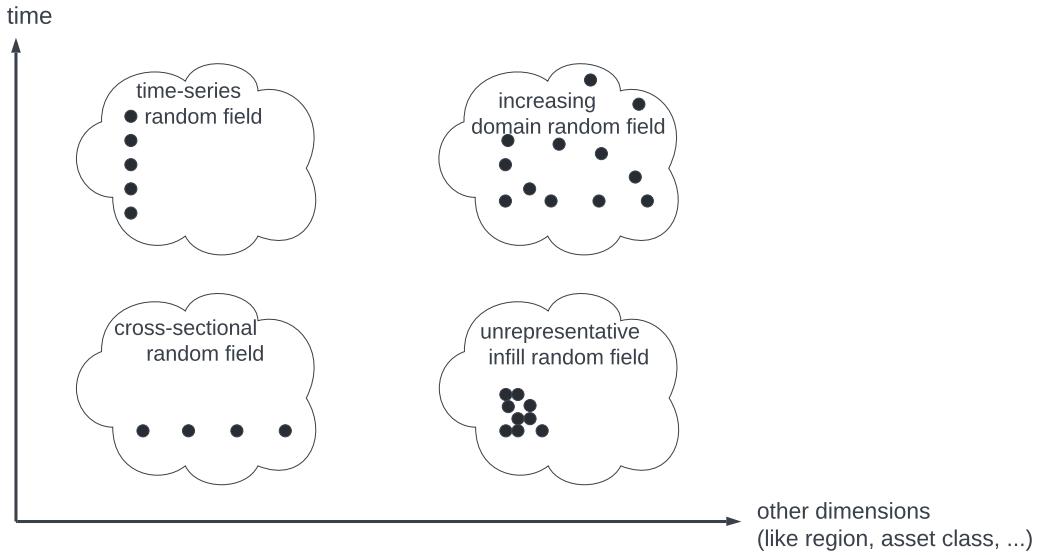


Figure 2: Visualization of generic random field types. Each black dot marks a different observation i of the cash flow data. Importantly, the time axis does **not** correspond to the index t in $CF_{i,t}$ (rather to vintage years v_i). Comparing the four choices, we want to avoid an infill random field but prefer our data to constitute an increasing domain random field. The infill random field is even asymptotically “too clustered” or better “too unrepresentative” to allow for meaningful estimation and inference. The time-series and cross-sectional random fields correspond to the standard cases in the literature but could turn out too restrictive for a general approach. By smart design (like portfolio formation), we often can map an increasing domain random field to simpler time-series or cross-sectional versions.

expose our SDF to many distinct covariate realizations (economic conditions), we also want the number of vintages to increase asymptotically.

Assumption 4. *Vintage year asymptotics:*

1. *The number of vintage years $V \rightarrow \infty$ as $n \rightarrow \infty$.*
2. *The number of funds per vintage year is bounded by some positive constant.⁴*
3. *The maximum fund lifetime is also bounded by a positive constant.*
4. *The economic distance between fund i and j is measured by the vintage year difference $d_{i,j} = |v_i - v_j| + \rho_0 \mathbf{1}_{i \neq j}$ with minimum distance $\rho_0 > 0$. Note that two distinct funds from the same vintage year have distance $d_{i,j} = \rho_0$, reflecting*

⁴This assumption is naturally satisfied when forming vintage-year portfolios (case (ii) of Assumption 3, with $n_v = 1$). For the individual fund case (i), it ensures that the SHAC estimator remains consistent and that standard errors scale with the number of vintage years rather than the total fund count. Relaxing this assumption would require cluster-robust inference at the vintage-year level.

our assumption that within-vintage-year dependence is captured by the SHAC kernel at the smallest lag.

In terms of the spatial estimation literature, this assumption postulates increasing domain asymptotics and rules out so-called infill asymptotics (cf. Figure 2). The minimum distance term ρ_0 is a means to ensure these increasing domain asymptotics (Jenish and Prucha, 2012, Assumption 1). Infill asymptotics corresponds to the approach of Driessen et al. (2012), who fix the number of portfolios N and let the number of underlying funds per portfolio $n_i \rightarrow \infty$. While this yields consistency for a given set of portfolios, it fundamentally conditions on the observed macro environment: with fixed N , no new vintage years and hence no new factor realizations enter the estimation. In contrast, our increasing domain framework ($V \rightarrow \infty$) permits inference that accounts for cross-vintage variation in economic conditions, which is necessary for the unconditional identification of SDF parameters (see Section 2.1). Finally, a further advantage of the *nonlinear least squares* formulation in Equation 18 is that it avoids the well-known complications arising when the number of GMM moment conditions grows with sample size (Han and Phillips, 2006; Newey and Windmeijer, 2009).

2.3.2 Law of large numbers

Under Assumption 2 (NPV-only discounting), Proposition 1 guarantees $\mathbb{E}[\bar{\epsilon}_i(\theta_0)] = 0$. For the general case where \mathcal{T}_i includes future value dates, a residual bias may persist (see Section 2.1). In both cases, the LLN ensures that the sample objective $Q_n(\theta)$ converges uniformly to its population counterpart. To approach this expected value in $\mathbb{E}[\bar{\epsilon}_i(\theta_0)]$, we rely on a spatial law of large numbers instead of applying a time-series law of large numbers. Here, we want to explicitly acknowledge the statistical dependence of pricing errors from adjacent vintage years.

Assumption 5. *Primitive conditions for uniform convergence (cf. Jenish and Prucha, 2009, Theorem 6):*

1. Near-epoch dependence (NED): *For each $\theta \in \Theta$, the random field $\{L(\bar{\epsilon}_i(\theta))\}_{i=1,\dots,n}$ is NED of size $-a$ (with $a > 0$) on an α -mixing random field $\{\eta_v\}_{v \geq 1}$ indexed by vintage years, i.e., $\|L(\bar{\epsilon}_i(\theta)) - \mathbb{E}[L(\bar{\epsilon}_i(\theta)) | \sigma(\eta_v : |v - v_i| \leq m)]\|_p \leq c_i \cdot \nu(m)$ with $\sup_i c_i < \infty$ and $\nu(m) = O(m^{-a})$.*
2. Mixing decay: *The base random field $\{\eta_v\}$ is α -mixing with mixing coefficients $\alpha(d) = O(d^{-\gamma})$ for some $\gamma > 2(p-1)/(p-2)$ where $p > 2$.*
3. Moment conditions: $\sup_{\theta \in \Theta} \mathbb{E}[|L(\bar{\epsilon}_i(\theta))|^p] < \infty$ for some $p > 2$, uniformly in i .
4. Stochastic equicontinuity: *The family $\{L(\bar{\epsilon}_i(\cdot)) : i = 1, \dots, n\}$ is stochastically equicontinuous in θ , which is satisfied if L is Lipschitz and $\bar{\epsilon}_i(\theta)$ is Lipschitz in θ uniformly in i .*

Under Assumptions 3–5, the ULLN $\sup_{\theta \in \Theta} |Q_n(\theta) - \mathbb{E}[Q_n(\theta)]| \xrightarrow{P} 0$ follows from (Jenish and Prucha, 2009, Theorem 6). Economically, the NED condition means that the pricing error of a fund in vintage v can be well-approximated using only

information from nearby vintages $\{v - m, \dots, v + m\}$, with approximation quality improving in m , which is natural since cash flow overlap decays with vintage year distance (cf. Assumption 4). In practice, we treat funds with distance $d_{i,j} > D$ as approximately independent for the purpose of covariance estimation (cf. the bandwidth parameter $b_n = D$ in Section 2.4). Since nominal fund sizes tend to grow over time, unweighted pricing errors $\bar{\epsilon}_i$ may be non-stationary; the weighting factor w_i from Equation 17 (e.g., scaled by invested capital) normalizes the pricing errors so that the moment condition (3) is satisfied.

2.3.3 Consistency

The estimator $\hat{\theta}$ shall converge in probability to the true parameter value θ_0 as the number of distinct vintage years goes to infinity.

Assumption 6. (*Identification*) *The population objective $Q_0(\theta)$ has a unique maximizer at θ_0 , i.e., $Q_0(\theta) < Q_0(\theta_0)$ for all $\theta \neq \theta_0$.*

Remark 2. *For $L(x) = x^2$, this requires that no alternative parameter vector $\theta \neq \theta_0$ can simultaneously achieve zero bias ($\mathbb{E}[\bar{\epsilon}(\theta)] = 0$) and equal variance ($\text{Var}[\bar{\epsilon}(\theta)] = \text{Var}[\bar{\epsilon}(\theta_0)]$). This is economically natural: mis-specifying the SDF factor loadings introduces systematic pricing errors whose variance exceeds the irreducible noise at θ_0 . The condition is analogous to the standard rank condition in linear GMM.*

Proposition 2 (Consistency). *Under Assumptions 3–6, $\hat{\theta} \xrightarrow{P} \theta_0$ as $n \rightarrow \infty$.*

Proof. We verify the four conditions of (Newey and McFadden, 1994, Theorem 2.1) for the extremum estimator $\hat{\theta} = \arg \max_{\theta \in \Theta} Q_n(\theta)$ with population objective $Q_0(\theta) := \mathbb{E}[Q_n(\theta)] = -\mathbb{E}[L(\bar{\epsilon}_i(\theta))]$.

1. *Identification.* Since $L \geq 0$, we have $Q_0(\theta) \leq 0$ for all θ . For the case $L(x) = x^2$,

$$Q_0(\theta) = -\text{Var}[\bar{\epsilon}(\theta)] - (\mathbb{E}[\bar{\epsilon}(\theta)])^2.$$

Under Assumption 2, Proposition 1 yields $\mathbb{E}[\bar{\epsilon}(\theta_0)] = 0$, so the bias term vanishes at θ_0 . For any $\theta \neq \theta_0$, at least one of (a) the bias term $(\mathbb{E}[\bar{\epsilon}(\theta)])^2 > 0$ or (b) the variance term $\text{Var}[\bar{\epsilon}(\theta)] > \text{Var}[\bar{\epsilon}(\theta_0)]$ must hold, ensuring $Q_0(\theta) < Q_0(\theta_0)$. This is the standard identification condition requiring that the true parameter is the unique minimizer of the expected loss.

2. *Compactness.* The parameter space Θ is compact by economic reasoning: factor loadings and intercepts are bounded by plausible ranges implied by the risk–return tradeoff.
3. *Continuity.* Continuity of $Q_0(\theta)$ follows from the continuous dependence of $\Psi_{\tau,t}(\theta)$ on θ and the dominated convergence theorem, which is a standard regularity condition.
4. *Uniform convergence.* By Assumption 5, $\sup_{\theta \in \Theta} |Q_n(\theta) - Q_0(\theta)| \xrightarrow{P} 0$, which directly satisfies the ULLN requirement of (Newey and McFadden, 1994, Section 2.1).

□

2.3.4 Central limit theorem

To assess the large-sample significance of our parameter estimates in Subsection 2.4, we establish the asymptotic distribution of $\hat{\theta}$.

Proposition 3 (Asymptotic Normality). *Under the conditions of Proposition 2, suppose additionally that (i) θ_0 is interior to Θ , (ii) $Q_0(\theta)$ is twice continuously differentiable at θ_0 with nonsingular Hessian $H := \nabla_{\theta\theta}Q_0(\theta_0)$, and (iii) a spatial CLT holds for the score. Then*

$$\sqrt{n}(\hat{\theta} - \theta_0) \xrightarrow{d} \mathcal{N}(0, \Sigma) \quad \text{as } n \rightarrow \infty,$$

with sandwich covariance matrix $\Sigma = H^{-1}\Lambda H^{-1}$, where $\Lambda := \lim_{n \rightarrow \infty} n \cdot \text{Var}[\nabla_{\theta}Q_n(\theta_0)]$ is the long-run covariance matrix of the score.

Proof. The result follows from the general asymptotic normality theorem for extremum estimators (Newey and McFadden, 1994, Theorem 3.1). The key non-standard ingredient is condition (iii): a CLT for the spatially dependent score $\nabla_{\theta}Q_n(\theta_0)$. Under the spatial near-epoch dependence structure posited in Assumption 5 and the vintage year asymptotics of Assumption 4, this follows from (Jenish and Prucha, 2012, Theorem 4), which extends the Pötscher and Prucha (1997) framework to spatial random fields. The remaining conditions of (Newey and McFadden, 1994, Theorem 3.1), i.e., interior θ_0 , smooth Q_0 , and stochastic equicontinuity of the score, are standard regularity conditions satisfied by the continuous dependence of $\Psi_{\tau,t}(\theta)$ on θ . \square

Remark 3. *The spatial CLT of (Jenish and Prucha, 2012, Theorem 4) requires the same NED and mixing conditions as Assumption 5, supplemented by a Lindeberg-type condition on the score contributions $\nabla_{\theta}L(\bar{\epsilon}_i)$. The Lindeberg condition is implied by the uniform moment bound in Assumption 5(3) applied to the score.*

2.4 Large sample inference

From Proposition 3, the asymptotic covariance matrix takes the sandwich form $\Sigma = H^{-1}\Lambda H^{-1}$ with H and Λ as defined therein. We now describe how to construct consistent estimators \hat{H} and $\hat{\Lambda}$.

We define the corresponding finite sample estimators analogously to Pötscher and Prucha (1997, Chapters 12, 13.1), and numerically approximate the first and second partial derivatives by finite differences⁶. Specifically, we use the following central difference approximations (with "small" δ) (Eu, 2017, Algorithm 2):

$$f_x(x, y) \approx \frac{f(x + \delta, y) - f(x - \delta, y)}{2\delta}$$

$$f_{xx}(x, y) \approx \frac{f(x + \delta, y) + f(x - \delta, y) - 2f(x, y)}{\delta^2}$$

⁵Since H is the Hessian of a scalar objective, it is symmetric, so $H^{-1} = (H^{-1})^\top$.

⁶As an alternative to finite differences the widespread Broyden-Fletcher-Goldfarb-Shanno (BFGS) algorithm can be applied to approximate the Hessian (Nocedal and Wright, 2006, Section 6.1).

$$f_{xy}(x, y) \approx \frac{f(x + \delta, y + \delta) + f(x - \delta, y - \delta) - f(x + \delta, y - \delta) - f(x - \delta, y + \delta)}{4\delta^2}$$

The Hessian term \hat{H} is relatively straightforward

$$\hat{H} = -\frac{1}{n} \sum_{i=1}^n \nabla_{\theta\theta} L(\bar{\epsilon}_i)$$

The consistency $\hat{H} \xrightarrow{P} H$ is implied by the ULLN (Assumption 5) applied to the second derivatives. Due to the spatial near-epoch dependence, the involved and computationally expensive part is to consistently estimate $\hat{\Lambda}$ by a Spatial Heteroskedasticity and Autocorrelation Consistent (SHAC) covariance matrix estimator (Kim and Sun, 2011, Equation 2)

$$\hat{\Lambda} = \frac{1}{n} \sum_{i=1}^n \sum_{j=1}^n k_{i,j} \left[\nabla_{\theta} L(\bar{\epsilon}_i) (\nabla_{\theta} L(\bar{\epsilon}_j))^{\top} \right] \quad (19)$$

Since $\nabla_{\theta} Q_n(\theta) = -\frac{1}{n} \sum_i \nabla_{\theta} L(\bar{\epsilon}_i)$, the score contributions are $s_i(\theta) := \nabla_{\theta} L(\bar{\epsilon}_i)$, and $\hat{\Lambda}$ equivalently estimates $n \cdot \text{Var}[\nabla_{\theta} Q_n]$.

We define the kernel weight k as

$$k_{i,j} \equiv K\left(\frac{d_{i,j}}{b_n}\right)$$

with kernel function $K : \mathbb{R} \rightarrow [0, 1]$ that satisfies $K(0) = 1$, $K(x) = K(-x)$, $\int_{-\infty}^{\infty} K^2(x) dx < \infty$, and $K(\cdot)$ continuous at zero and at all but a finite number of other points. A common choice is the Bartlett kernel $K_{BT}(x) = \max(0, 1 - |x|)$; see equation 2.7 in Andrews (1991) for other popular kernel choices. This means absolute vintage year differences larger than the bandwidth (or truncation) parameter $b_n = D$ are considered independent and are thus excluded from the $\hat{\Lambda}$ estimation formula.

In large samples, the vector of parameter standard errors can thus be estimated by

$$\text{SE}(\hat{\theta}) = \sqrt{\text{diag} \left[n^{-\frac{1}{2}} \hat{H}^{-1} \hat{\Lambda} \hat{H}^{-1} (n^{-\frac{1}{2}})^{\top} \right]} = \sqrt{\text{diag} \left[\hat{H}^{-1} \hat{\Lambda} \hat{H}^{-1} \right] \cdot \frac{1}{n}} \quad (20)$$

The Wald test statistic for linear hypotheses $\mathcal{H}_0 : R\theta = r$ and $\mathcal{H}_1 : R\theta \neq r$ is constructed as (where \mathcal{H}_0 and \mathcal{H}_1 are hypotheses and not Hessian terms H)

$$W = (R\hat{\theta} - r)^{\top} \left[R \frac{\hat{H}^{-1} \hat{\Lambda} \hat{H}^{-1}}{n} R^{\top} \right]^{-1} (R\hat{\theta} - r) \stackrel{H_0}{\sim} \chi_q^2$$

where $\hat{\theta}$ is the $p \times 1$ parameter vector, R is a $q \times p$ matrix, and r is a $q \times 1$ vector. Usually, we select R as $p \times p$ identity matrix, and r as $p \times 1$ vector (e.g., of zeros). Under the null hypothesis, W is chi-squared distributed with q degrees of freedom. As large values of W indicate the rejection of H_0 , the corresponding p-value is calculated as $1 - F_{\chi_q^2}(W)$ where $F_{\chi_q^2}$ is the cumulative distribution function of a chi-squared random variable with q degrees of freedom.

However, given the limited amount of available private equity data (typically the oldest vintages start in the 1980s), asymptotic characterizations of Σ and $\text{SE}(\hat{\theta})$ are of limited importance. In empirical applications, the small sample behavior of an estimation method for private equity data is more relevant than its asymptotic theory. Moreover, the standard asymptotic distribution associated with an estimator is generally not valid for post-model-selection inference, i.e., if a model selection procedure is applied to find the best model from a collection of competitors (Leeb and Pötscher, 2005).

2.5 Comparison to similar estimators

Our Least-Mean-Distance (LMD) estimator introduced in Section 2.2 belongs to the class of semiparametric nonlinear M-estimators as defined in Pötscher and Prucha (1997) which are extremum estimators. To gain more flexibility and avoid unneeded complexity, we intentionally opt against the most prominent semiparametric nonlinear M-estimator framework, i.e., classical time-series Generalized Method of Moments (GMM) (Hansen, 1982, 2012). A classical GMM approach requires the construction of stationary, ergodic time-series of moment conditions that are used to empirically estimate the expected value of pricing errors in Equation 16. The stationarity requirement of classical time-series GMM limits (i) more elaborate weighting schemes for w , like fund-size weighting, and (ii) the usage of fund cash flows from non-realized vintages.

The flexibility of our framework directly embeds the two most prominent asset pricing approaches in the private equity literature as special cases. We formalize this relationship below to demonstrate that our proposed methodology is a true generalization of both the cross-sectional approach of Driesssen et al. (2012) and the time-series approach of Korteweg and Nagel (2016).

Proposition 4 (Generalization of existing estimators). *The LMD estimator defined in Equation 18 generalizes the following two estimators:*

1. **Driesssen et al. (2012):** *By restricting the set of discounting dates to fund inception ($\mathcal{T}_i = \{\min_j t_{i,j}^{\text{Inv}}\}$), applying equal weights ($w_i = 1$), and choosing a quadratic loss function ($L(x) = x^2$), the LMD minimization objective becomes identical to their cross-sectional Generalized Method of Moments (GMM) objective function.*
2. **Korteweg and Nagel (2016):** *By interchanging the order of the summation and the loss function evaluation in the LMD minimization objective (i.e., evaluating $L\left(\frac{1}{n} \sum_{i=1}^n \bar{\epsilon}_i\right)$ instead of $\frac{1}{n} \sum_{i=1}^n L(\bar{\epsilon}_i)$), while defining the loss function as the standard quadratic GMM objective ($L(x) = x^\top W x$) with weighting matrix W and pooling all funds into a single cross-sectional unit ($n = 1$), our framework yields their time-series GMM estimator.*

2.5.1 Comparison to Driesssen et al. (2012)

The Driesssen et al. (2012) approach is most closely related to our methodology. One important difference is that we select a simpler and more flexible LMD estimator instead of a cross-sectional GMM approach. In our view, the choice of the

more complex cross-sectional GMM just introduces conceptual complexity without changing the estimating equations since the underlying formulas are basically the same as for our LMD estimator⁷.

Proof of Proposition 4 (Part 1). Applying the restrictions $\mathcal{T}_i = \{\min_j t_{i,j}^{\text{Inv}}\}$ and $w_i = 1$ to our average pricing error from Equation 17 yields exactly the inception-based net present value:

$$\bar{\epsilon}_i = \sum_{t=0}^T \frac{\Psi_t}{\Psi_{\min_j t_{i,j}^{\text{Inv}}}} \cdot CF_{i,t} = PV_i$$

Assuming equal weights for all funds leads to the minimization formulation of the LMD estimator in Equation 18 with quadratic loss $L(x) = x^2$:

$$\hat{\theta} = \arg \min_{\theta} \frac{1}{n} \sum_{i=1}^n (PV_i)^2 \quad (21)$$

This minimization objective is equivalent to the inner product of the pricing error vector with itself, which precisely recovers Equation 3 in Driessen et al. (2012): $\min_{\theta} e'e$, where e is an $n \times 1$ vector containing the pricing errors of the n cross-sectional units. \square

Consequently, if someone accepts the assumptions from Subsection 2.3, our large sample inference framework from Subsection 2.4 applies to their case without any significant modification.

As a first limitation, they have to regard vintage-year portfolios as their cross-sectional units; we can also use individual funds. In this context, Driessen et al. (2012) state that “to identify β , it is essential that the different FoFs are exposed to different market returns.” We note that their estimator is valid as a cross-sectional approach: identification requires only that the cross-section of funds provides sufficient variation in factor exposure at a given market return realization⁸. That said, having FoFs exposed to different market environments (i.e., different vintage years) does improve statistical power and is the basis for our increasing domain asymptotics.

Second, the Driessen et al. (2012) asymptotic theory is based on infill asymptotics: they fix the number of vintage year portfolios n and let the number of underlying funds per portfolio $n_i \rightarrow \infty$. As discussed in Section 2.3, this conditions inference on the observed macro environment since no new vintage years enter the estimation. Moreover, since each portfolio constitutes a moment condition from the GMM perspective, standard GMM theory requires n to remain fixed (Han and Phillips, 2006; Newey and Windmeijer, 2009). In contrast, our increasing domain framework lets the number of vintage years $V \rightarrow \infty$ while bounding the number of funds per vintage, enabling inference that reflects cross-vintage variation in economic conditions.

Further, Driessen et al. (2012) discount all fund cash flows exclusively to the fund inception date, corresponding to $\mathcal{T}_i = \{\min_j t_{i,j}^{\text{Inv}}\}$ in our notation. Under

⁷The formulas are this similar because Driessen et al. (2012) use the identity matrix as GMM weighting matrix and skip the second GMM step.

⁸In a classic cross-sectional regression, we only have one market return realization.

Assumption 2, this is also the only choice that yields unconditionally unbiased estimates. Our framework generalizes this by allowing \mathcal{T}_i to contain additional discounting dates, which we use to empirically study the bias–variance tradeoff analyzed in Section 2.1. Finally, Driessen et al. (2012) apply simple cross-sectional bootstrapping to obtain standard errors; in contrast, in Subsection 3.2, we use a cross-validation technique that is adapted to the near-epoch dependence of the PE fund data.

2.5.2 Comparison to Korteweg and Nagel (2016)

Korteweg and Nagel (2016) were the first to employ a spatial framework for asymptotic inference on private equity fund data. They measure the economic distance between funds by the degree of cash flow overlap and use the spatial HAC estimator of Conley (1999) to construct robust standard errors. Our spatial HAC framework similarly accounts for cross-vintage dependence but draws on the near-epoch dependent theory of Pötscher and Prucha (1997); Kim and Sun (2011); Jenish and Prucha (2012). A key methodological difference is that Korteweg and Nagel (2016) estimate the SDF on synthetic public market portfolios using a time-series GMM estimator and then apply the resulting model to evaluate PE cash flows, whereas we estimate the SDF directly from PE fund cash flows. Korteweg and Nagel (2016) choose this procedure since “tests that use the ex post equity premium as a performance benchmark should be more powerful than tests that use the ex ante equity premium.”

Proof of Proposition 4 (Part 2). The time-series GMM objective is formed by pooling the expected values of the moment conditions. If we treat each of the n cross-sectional units equivalently as a moment condition and compute the average pricing error across the panel, we obtain the m -dimensional sample moment $\frac{1}{n} \sum_{i=1}^n \bar{\epsilon}_i \in \mathbb{R}^m$. The standard GMM estimator minimizes the quadratic form of these moments weighted by a matrix W . Formally, defining $L(x) = x^\top W x$, the time-series GMM objective evaluated on the pooled sample moments becomes:

$$Q_{GMM}(\theta) = L\left(\frac{1}{n} \sum_{i=1}^n \bar{\epsilon}_i\right) = \left(\frac{1}{n} \sum_{i=1}^n \bar{\epsilon}_i\right)^\top W \left(\frac{1}{n} \sum_{i=1}^n \bar{\epsilon}_i\right) \quad (22)$$

Setting $W = I$ corresponds exactly to Equation 18 in Korteweg and Nagel (2016). Therefore, the minimization formulation of our LMD objective in Equation 18 (i.e., $\frac{1}{n} \sum_{i=1}^n L(\bar{\epsilon}_i)$) differs from Korteweg and Nagel (2016) solely by whether the loss function $L(\cdot)$ is evaluated before or after pooling the cross-sectional errors. \square

Time-series GMM estimators inherently bear the risk of under-identification if the corresponding time-series is constructed by pooling all fund cash flows from a given fund type. When PE fund cash flows are pooled into a single moment condition as in Equation 22 with $m = 1$, the resulting estimator has just one moment condition, which is insufficient to identify multi-parameter models⁹. Note that Korteweg and Nagel (2016) avoid this issue by estimating the SDF from public market

⁹In contrast, our estimator corresponds to the opposite edge case with asymptotically an infinite number of LMD “moment conditions” (units to price) as we let $n \rightarrow \infty$.

data rather than PE cash flows; the under-identification concern arises only if one applies their framework directly to PE data, as we do in our estimator. To counter under-identification, additional characteristic-based fund portfolios could be formed to increase the number of moment conditions per fund type; also, random portfolios combined with bootstrapping could make sense. Yet, Korteweg and Nagel (2016) take another approach and introduce the concept of Generalized Public Market Equivalent (GPME), which elegantly avoids the under-identification issue. Firstly, a public market SDF model is estimated by pricing public trading strategies that shall replicate PE funds instead of directly pricing the observed PE fund cash flows. Only in a second step are these models applied to evaluate private equity fund cash flows.

In summary, our LMD estimator generalizes Driessen et al. (2012) by (a) permitting flexible discounting dates via \mathcal{T}_i , (b) employing increasing domain rather than infill asymptotics, and (c) providing SHAC-based inference. Relative to Korteweg and Nagel (2016), we directly price PE cash flows rather than pricing public replicating strategies, which avoids potential misspecification of the public SDF but requires stronger identification assumptions from the PE data alone. Table 2 summarizes the most prominent distinctions between the three approaches.

	Driessen et al. (2012)	Korteweg and Nagel (2016)	Our approach
Estimator type	Cross-sectional NLS (one-step GMM)	Time-series GMM (public SDF)	Nonlinear Least-Mean-Distance
Cash flows priced	PE fund cash flows	Public replicating strategies	PE fund cash flows
Discounting dates	Fund inception only	Fund inception only	Flexible via \mathcal{T}_i
Asymptotics	Infill ($n_i \rightarrow \infty$, N fixed)	$V \rightarrow \infty$	Increasing domain ($V \rightarrow \infty$, n_v bounded)
Inference	Bootstrap	Spatial HAC	Spatial HAC
Cross-sectional unit	Vintage year portfolio	Single fund	Both (tested empirically)
SDF specification	Linear factor model	Exponentially affine	Both (tested empirically)

Table 2: Comparison of SDF estimation frameworks for private equity.

3 Empirical application

3.1 Data

We use the Preqin cash flow data set as of 14th April 2022 which is well known in the academic private equity literature (Harris et al., 2014; Brown et al., 2015; Korteweg and Nagel, 2016; Ang et al., 2018). To keep the empirical analysis clear and concise, we filter the dataset for Private Equity (PE) funds.

The Private Equity sample contains 2745 distinct funds spreading over 35 vintage years. The vintage year distribution is as follows: 1983 (1), 1985 (2), 1986 (4), 1987 (5), 1988 (7), 1990 (2), 1991 (4), 1992 (9), 1993 (13), 1994 (20), 1995 (16), 1996 (20), 1997 (24), 1998 (59), 1999 (46), 2000 (72), 2001 (51), 2002 (41), 2003 (45), 2004 (64), 2005 (117), 2006 (154), 2007 (176), 2008 (183), 2009 (83), 2010 (92), 2011 (146), 2012 (136), 2013 (142), 2014 (182), 2015 (152), 2016 (186), 2017 (143), 2018 (157), 2019 (190). The region distribution is as follows: Africa (20), Asia (150), Australasia (39), Europe (681), Latin America & Caribbean (33), Middle East (14), North America (1807). The strategy distribution is as follows: Balanced (69), Buyout (1354), Co-Investment (87), Co-Investment Multi-Manager (59), Direct Secondaries (30), Fund of Funds (643), Growth (331), Secondaries (152), Turnaround (19).

For this subset of Private Equity funds, we extract all (i) equal-weighted and (ii) fund-size-weighted cash flow series. For non-liquidated funds, we treat the latest net asset value as final cash flow. We explicitly refrain from excluding the most recent vintage years. Thus, the minimum vintage year is 1983 (just for PE), and the maximum is 2019 since we filtered out funds with younger vintage.

The public market factors that enter our SDF draw on the US data set of the recently popularized q^5 investment factor model sourced from <http://global-q.org/factors.html> (Hou et al., 2015, 2020). Their five-factor model includes the market excess return (MKT), a size factor (ME), an investment factor (IA), a return on equity factor (ROE), and an expected growth factor (EG).

3.2 Model and estimator specifications

We test a simple linear SDF model similar to Driessen et al. (2012)

$$\Psi_{\tau,t}^{\text{SL}}(\theta) = \frac{\prod_{h=0}^{\tau} (1 + \alpha + r_{f,h} + \sum_m \beta_m F_{m,h})}{\prod_{h=0}^t (1 + \alpha + r_{f,h} + \sum_m \beta_m F_{m,h})} \quad (23)$$

and an exponential affine SDF model adapted from Korteweg and Nagel (2016)

$$\Psi_{\tau,t}^{\text{EA}}(\theta) = \exp \left[\sum_{h=0}^{\tau} X_h \sum_{h=0}^t -X_h \right] \quad (24)$$

with

$$X_h = \alpha + \log(1 + r_{f,h}) + \sum_{m \in M} \beta_m \cdot \log(1 + F_{m,h})$$

with (arithmetic) risk-free return $r_f = R_{rf} - 1$, (arithmetic) zero-net-investment portfolio returns F_m , and parameter vector $\theta = (\alpha, \beta)$. To avoid overfitting, we just test six simple SDF models that contain {MKT} alone or {MKT} plus {ME or IA or ROE or EG or Alpha}. In Equation 18, we use the quadratic loss function $L(x) = x^2$.

To assess the parameter significance, we compute the asymptotic standard errors as outlined in Subsection 2.4. For the Bartlett kernel's bandwidth $b_n = D$ we select 12 years, i.e., funds with absolute vintage year differences larger than 12 years are assumed to be independent.

Additionally, we want to test the small-sample parameter significance and the out-of-sample performance of our SDF models. To account for the dependency between funds from adjacent vintage years caused by overlapping fund cash flows, we draw on *hv*-block cross-validation (Racine, 2000). Therefore, we form three partitions for several vintage year groups. As larger validation sets are preferred for model selection, the validation set (*v*-block) always contains funds of three neighboring vintage years (e.g., 2000, 2001, 2002). To reduce the dependency between training and validation set, we remove all funds from three-year-adjacent vintage years, i.e., the *h*-block (e.g., 1997, 1998, 1999, 2003, 2004, 2005). Funds from the remaining vintage years enter the training set and are thus used for model estimation (e.g., 1985-1996, 2006-2019). We apply ten-fold cross validation using the ten validation sets described in Table 3. This means we replace the bootstrap standard error calculation of Driesssen et al. (2012) by *hv*-block cross-validation. Since the finite sample dimensionality of PE datasets often causes asymptotic standard errors to be unreliable large and volatile, *hv*-block cross validation is necessary because it (i) accounts for near-epoch-dependence without relying on asymptotics, (ii) focuses directly on the out-of-sample performance of the SDF models, and (iii) is computationally cheaper.

Remark 4 (Definition of Horizon). *Throughout the empirical section, we use the term Horizon H to denote the maximum compounding horizon in years used in the set of discounting dates \mathcal{T} . Formally, assuming a contiguous set of monthly discounting dates $\mathcal{T} = \{1, 2, \dots, m_{\max}\}$ (where m_{\max} is the maximum month), the Horizon is defined as $H = m_{\max}/12$. Consequently, estimating a model with a Horizon H implies that the estimator minimizes the average pricing error over the set of discounting dates $\mathcal{T}_H = \{1, \dots, 12 \cdot H\}$, formally denoted as $\bar{\epsilon}_i^{(H)}$.*

Remark 5 (Alpha Bounds). *In our optimization procedure, we limit the Alpha factor α to the range of $[-2\%, +2\%]$ per month to avoid the “exploding alpha” issue documented in (Driesssen et al., 2012, Footnote 8): “Note that one needs to put an upper bound on α when performing the optimization in equation (3). This is to exclude the possibility that α tends to infinity.” In this degenerate scenario, an infinite alpha can almost perfectly price any cash flow stream, rendering the beta risk factors irrelevant.*

3.3 Simulation study

Our Monte Carlo experiments examine the bias and variance of our estimation methodology in finite samples.¹⁰ Specifically, we investigate five key dimensions of the estimation framework: (i) the finite-sample efficiency gains of forming vintage-year portfolios versus using individual funds; (ii) the tradeoff between expanding the time-series dimension (number of vintage years) and the cross-sectional dimension (funds per vintage); (iii) the behavior of multivariate identification under idiosyncratic noise using two-factor models; (iv) the robustness of the estimates to the

¹⁰As with any simulation study, it investigates the ability to identify the assumed data generating process rather than validate the corresponding SDF model itself.

training.before estimation	h -block.before remove	v -block validation	h -block.after remove	training.after estimation
start-1984	1985,1986,1987	1988,1989,1990	1991,1992,1993	1994-end
start-1987	1988,1989,1990	1991,1992,1993	1994,1995,1996	1997-end
start-1990	1991,1992,1993	1994,1995,1996	1997,1998,1999	2000-end
start-1993	1994,1995,1996	1997,1998,1999	2000,2001,2002	2003-end
start-1996	1997,1998,1999	2000,2001,2002	2003,2004,2005	2006-end
start-1999	2000,2001,2002	2003,2004,2005	2006,2007,2008	2009-end
start-2002	2003,2004,2005	2006,2007,2008	2009,2010,2011	2012-end
start-2005	2006,2007,2008	2009,2010,2011	2012,2013,2014	2015-end
start-2008	2009,2010,2011	2012,2013,2014	2015,2016,2017	2018-end
start-2011	2012,2013,2014	2015,2016,2017	2018,2019,2020	2021-end

Table 3: Partitions used for hv -block cross-validation.

functional specification of the SDF (simple linear vs. exponential affine); and (v) the optimal choice of the maximum compounding horizon (the set of discounting dates \mathcal{T}).

We use historical q -investment factors from 1986 to 2005 and simulate 20 funds for each of these 20 vintage years. This means we condition our simulation analysis on realized q -investment factors paths (see the Macro-Conditioning Kernel discussion in Section 2.1). Each fund contains 15 deals with equal investment amounts and exactly one divestment cash flow. Deals are entered within the first five years of the fund lifetime following a discrete uniform distribution and afterward held between one to ten years again uniformly distributed. The deal returns are generated by the simple linear or exponential affine SDF models described in Equations 23 and 24. In the base case, we just use the MKT factor with $\beta_{MKT} = 1$ and in each month, add a normal i.i.d. error term with monthly standard deviation $\sigma = 0.2$ (annualized variance $12 \cdot 0.2^2 = 0.48$, corresponding to an annualized volatility of $0.2\sqrt{12} \approx 69\%$) and zero mean. In the exponential affine case, we adjust the log-normally distributed error mean to zero by subtracting $0.5\sigma^2$. If a negative return exceeds -100%, the company defaults with a zero exit cash flow. In contrast, the error term in the simulations of Driessen et al. (2012) is more well-behaved as it follows a shifted lognormal distribution that, even with arbitrarily high error term variance, just allows for returns below say -99%, if the market return is close to its lower bound (see equation 9 in their online appendix; we specifically analyze the remaining divergence from their simulation design in Appendix A). In our base case, the set of discounting dates \mathcal{T} contains all months from the first cash flow to the maximum month 180. To assess our estimator's bias and variance, we simulate 1000 test scenarios for both vintage-year and single-fund portfolios.

Cross-sectional unit i : As presumed in Subsection 2.3, vintage year portfolio results appear to have lower bias and variance when compared to individual funds. For the simple linear SDF and 15-year horizon (maximum month 180), the mean and standard deviation of the coefficient estimate $\hat{\beta}_{MKT}$ is 0.966 (0.188) for the vintage year portfolio and 0.876 (0.353) for individual funds. The superiority of the portfolio

approach becomes even more apparent when considering the entire horizon curve in Figure 3. Estimates for individual funds are highly unstable, exhibiting severe upward bias at short horizons ($\hat{\beta}_{MKT} \approx 1.5$) and deteriorating to negative values at long horizons. In contrast, vintage year portfolios yield substantially more stable estimates that remain within a pertinent economic range around the true value of 1.0. Furthermore, the portfolio approach consistently achieves a lower standard deviation at every simulated horizon, confirming that grouping funds effectively mitigates idiosyncratic noise.

This finding has two important implications: On the one hand, vintage year portfolio formation can substantially decrease our estimator's bias and variance. On the other hand, it also dramatically reduces the number of cross-sectional units and consequentially impairs the importance of asymptotic results. These considerations may explain the choice of Korteweg and Nagel (2016) to use individual funds as cross-sectional units in their asymptotic SHAC framework to obtain smaller standard error estimates.

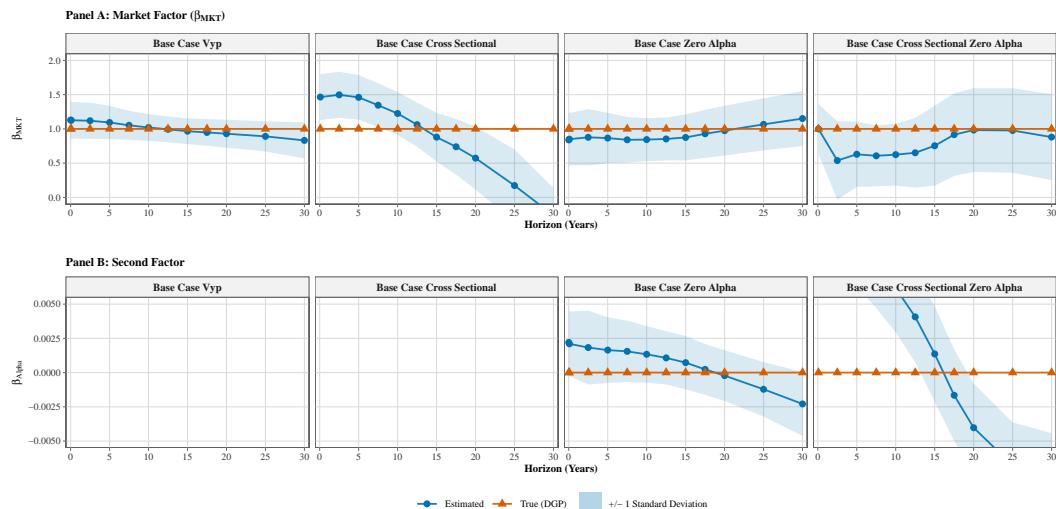


Figure 3: **Simulation** (β_{MKT} , α): Single-fund vs. vintage-year portfolios.

Varying vintages V and portfolio sizes n/V : A fundamental question in private equity panel data econometrics is whether precision is best improved by extending the time-series dimension (V , number of vintage years) or the cross-sectional dimension (n/V , number of funds per vintage portfolio). To assess this tradeoff, we compare five data configurations (depicted in Figures 4 and 5): (1) our base case ($V = 20$, $n/V = 40$); (2) an extended vintage sample with small portfolios ($V = 39$, $n/V = 10$); (3) an extended vintage sample with medium portfolios ($V = 39$, $n/V = 20$); (4) an early-vintage sub-sample ($V = 10$, years 1986–1995, $n/V = 20$); and (5) a late-vintage sub-sample ($V = 10$, years 1996–2005, $n/V = 20$).

Our simulation yields a clear prescriptive result for empirical practice: increasing the number of cross-sectional units per portfolio (n/V) is substantially more effective at suppressing estimator variance than adding more vintage years (V). For the

single-factor model at the 15-year horizon, transitioning from 10 funds to 20 funds per vintage over a 39-year period (scenarios 2 vs. 3) reduces the standard deviation of $\hat{\beta}_{MKT}$ by roughly 22%, from 0.218 to 0.169. In contrast, shrinking the time-series dimension from 39 years to just 10 years (while holding $n/V = 20$ constant) elevates the variance only moderately for the early-vintage sub-sample ($\sigma(\hat{\beta}_{MKT}) \approx 0.194$ for 1986–1995), though the late-vintage sub-sample (1996–2005) exhibits considerably higher variance ($\sigma(\hat{\beta}_{MKT}) \approx 0.432$), reflecting the sensitivity to the realized macro environment. This hierarchy of variance reduction is even more pronounced in the two-factor setting, where idiosyncratic cash flow noise severely impairs multivariate identification: increasing n/V aggressively dampens the spurious α variance. Consequently, our findings provide a compelling rationale for the finite-sample superiority of the Driesssen et al. (2012) infill asymptotics (fixing V , letting $n_v \rightarrow \infty$) over approaches that rely purely on time-series aggregation.

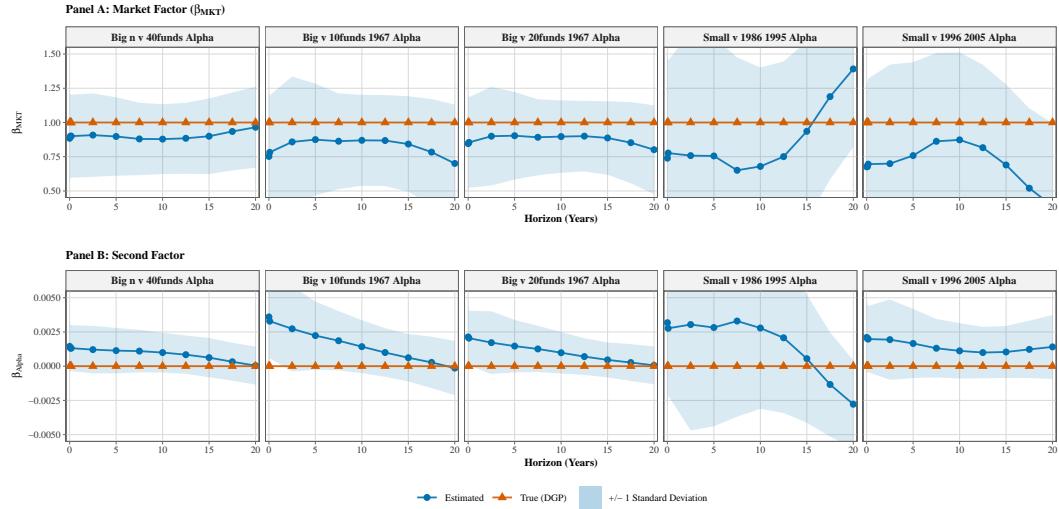


Figure 4: **Simulation (β_{MKT}, α):** Varying vintages V and portfolio sizes n/V .

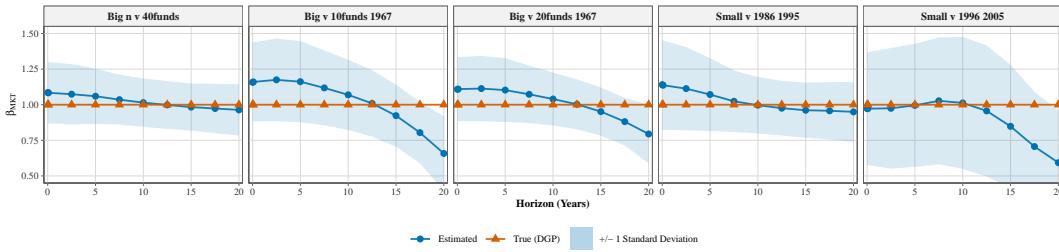


Figure 5: **Simulation (β_{MKT}):** Varying vintages V and portfolio sizes n/V .

SDF model Ψ : In Figure 6, we compare the simple linear SDF model against the exponential affine specification. For the single-factor model, the bias is minimized at the 12.5-year horizon (150 months), where both functional forms perform

almost identically and yield estimates extremely close to the true value of 1.0: the exponential affine yields a mean market beta of 0.999 (0.182), matching the simple linear's 0.994 (0.189). For the two-factor (MKT+Alpha) model at this same 12.5-year horizon, both functional forms again yield nearly identical parameter results, though with an attenuated market beta ($\hat{\beta}_{\text{MKT}} \approx 0.85$) and a minor spurious positive alpha ($\hat{\alpha} \approx 0.001$). In the two-factor case, the market beta bias is actually minimized later, around the 20-year horizon ($\hat{\beta}_{\text{MKT}} \approx 0.97$, $\hat{\alpha} \approx -0.0002$).

However, a salient difference emerges at extremely short compounding horizons. At a 1-month horizon (resembling an NPV-only discounting framework), the exponential affine formulation deteriorates much more severely in both models, manifesting a larger upward bias (e.g., in the single-factor case, $\hat{\beta}_{\text{MKT}} = 1.21$ vs. 1.13) and elevated variance (0.43 vs. 0.27) compared to the linear model. Overall, the perceived superiority of exponential affine SDFs (Gourieroux and Monfort, 2007; Bertholon et al., 2008) does not translate into empirical finite-sample advantages in our framework. The simple linear model actually appears slightly better behaved.

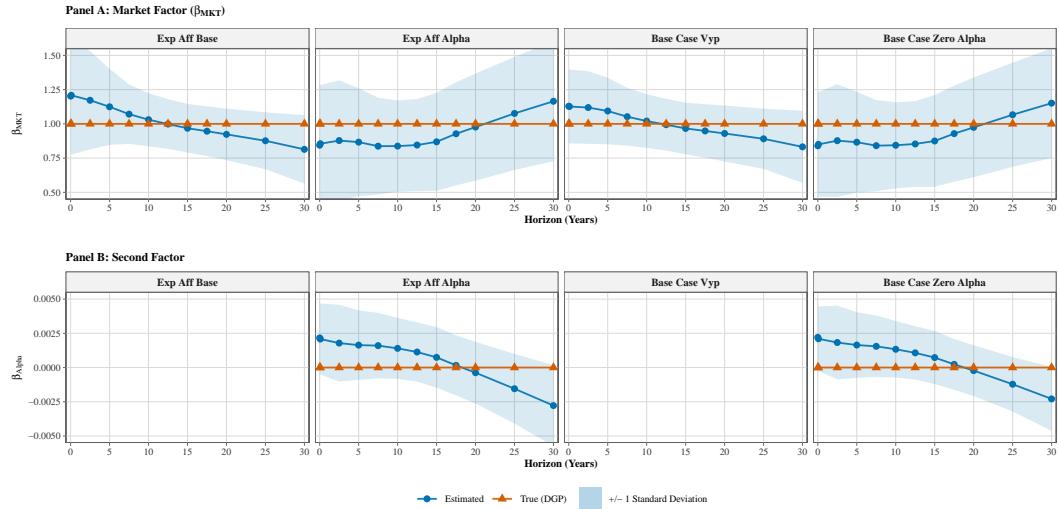


Figure 6: **Simulation ($\beta_{\text{MKT}}, \alpha$):** Exponential affine vs. linear models.

Two-factor models (q^5 factors): Figure 7 reports simulation results for two-factor models that pair MKT with each of the four non-market q^5 factors (ME, IA, ROE, EG), with true loadings $\beta_{\text{MKT}} = 1$ and $\beta_{\text{2nd}} = 0.5$. A common pattern emerges across all four specifications: the second-factor loading is substantially overestimated at short horizons and attenuated toward — or beyond — the true value at long horizons, while the market factor absorbs a compensating bias in the opposite direction. Estimation variance is markedly larger than in the single-factor case, reflecting the collinearity between correlated factors in a limited cross-section of vintage year portfolios. Table 4 summarizes the key estimates at the shortest horizon and at 15 years, which is approximately the bias-minimizing horizon for most models; we highlight the most notable factor-specific features below.

The ME model is the only specification in which $\hat{\beta}_{\text{2nd}}$ undergoes a sign change

(at roughly 17.5 years), declining from 0.86 to -0.02 over the full horizon range — a severe attenuation that would qualitatively alter economic conclusions. For the IA model, the second-factor bias is large at short horizons ($\hat{\beta}_{\text{IA}} \approx 1.06$ vs. truth 0.5) but both loadings converge close to their true values near the 17.5-year horizon ($\hat{\beta}_{\text{MKT}} = 0.97$, $\hat{\beta}_{\text{IA}} = 0.52$). The EG model exhibits the largest overall estimation variance ($\sigma(\hat{\beta}_{\text{EG}}) \approx 2.0$ at the shortest horizon) but achieves a reasonably accurate joint estimate around 15–17.5 years ($\hat{\beta}_{\text{MKT}} = 0.91\text{--}1.02$, $\hat{\beta}_{\text{EG}} = 0.47\text{--}0.56$). The ROE model is the exception: both loadings remain persistently biased at all horizons, with $\hat{\beta}_{\text{MKT}}$ never exceeding 0.91 and $\hat{\beta}_{\text{ROE}}$ never falling below 0.62. These results suggest that empirical researchers should exercise caution when interpreting two-factor SDF estimates for private equity, and that the bias-minimizing horizon is model-dependent, typically falling in the 12.5–17.5-year range.

Model	$\hat{\beta}_{\text{MKT}}$		$\hat{\beta}_{\text{2nd}}$		$\sigma(\hat{\beta}_{\text{MKT}})$		$\sigma(\hat{\beta}_{\text{2nd}})$	
	0 yr	15 yr	0 yr	15 yr	0 yr	15 yr	0 yr	15 yr
ME + MKT	1.15	0.97	0.86	0.59	0.28	0.21	0.46	0.40
IA + MKT	0.97	0.92	1.06	0.64	0.34	0.30	0.80	0.53
ROE + MKT	0.84	0.83	1.04	0.69	0.43	0.49	0.72	0.49
EG + MKT	0.87	0.91	1.04	0.56	0.65	0.45	2.00	0.31

True values: $\beta_{\text{MKT}} = 1$, $\beta_{\text{2nd}} = 0.5$.

Table 4: Two-factor simulation summary: mean estimates and standard deviations at the shortest (0 yr) and the 15-year horizon.

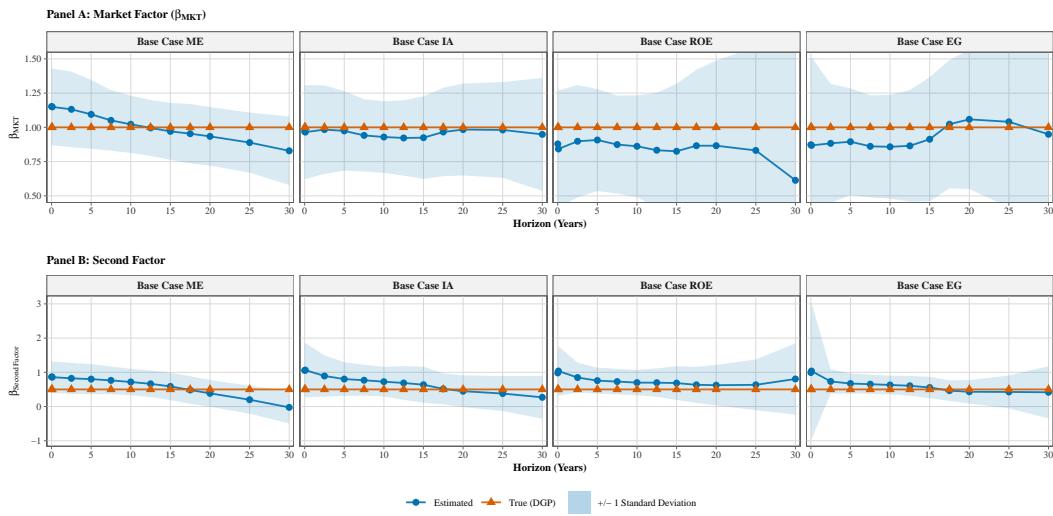


Figure 7: **Simulation (q^5 factors):** Two-factor models using VYP.

Size of set \mathcal{T} (Maximum Compounding Horizon): Our most critical simulation finding relates to the optimal choice of the maximum compounding horizon,

i.e., the size of the set of discounting dates \mathcal{T} . From a purely theoretical perspective (cf. Section 2.1), discounting all cash flows exclusively to the fund inception date, corresponding to a zero-month or NPV-only horizon in our notation, should yield unconditionally unbiased estimates under Assumption 2. This minimal choice for \mathcal{T} corresponds exactly to the established Driessen et al. (2012) framework.

However, our diverse simulation results reveal a puzzling disconnect between asymptotic theory and finite-sample performance: extending the set of discounting dates deep into the fund lifecycle drastically improves the estimator’s precision. Across almost all model specifications (see Figures 7 and 6), bias is minimized not at the theoretical 0-year horizon, but rather around a 12.5 to 15-year compounding horizon. At extremely short horizons, idiosyncratic cash flow noise interacts non-linearly with the estimation procedure, inducing severe upward bias in factor loadings. By averaging pricing errors over multiple future dates, effectively incorporating the Timing Risk Premium, our approach acts as a powerful finite-sample bias correction.

Interestingly, this “variance reduction technique” improves the mean estimate far more than it reduces standard errors. While estimator variance traces a shallow U-shape with a minimum near the 15-year horizon (see Table 4), this reduction is secondary to the striking stabilization of the parameter bias. We further analyze this puzzling result from a theoretical perspective in Appendix A, detailing the *bias–variance tradeoff* and the role of the Timing Risk Premium.

Simulation conclusion: To conclude, our simulation study rationalizes two key practices from the Driessen et al. (2012) paper: (i) vintage-year portfolio formation helps to improve estimator precision, and (ii) increasing the number of funds per vintage seems more effective in controlling estimator variance than increasing the number of vintages. Our second finding provides an empirical rationale for the choice of Driessen et al. (2012) to employ infill asymptotics that let the number of underlying funds per vintage tend to infinity. However, our correctly specified scenarios cannot support the assumption of Korteweg and Nagel (2016) that (iii) the exponential affine SDF is superior to the simple linear SDF in a multi-period framework; in fact, their finite-sample bias and variance profiles are remarkably similar. Moreover, our simulation study suggests that (iv) averaging pricing errors over multiple dates substantially reduces the severe short-horizon bias inherent to the original procedure of Driessen et al. (2012), which discounts all cash flows exclusively to the fund inception date. Ultimately, selecting a set of compounding dates \mathcal{T} that spans a significant portion of the fund lifecycle appears to decrease finite-sample bias (and to a lesser extent, variance) more effectively than all other measures combined. Appendix A explains the theoretical foundation of this observation, illustrating how our biased estimator (cf. Eq. 10) achieves superior finite-sample properties by exploiting the bias–variance tradeoff.

3.4 Empirical results

Following the conclusions from the simulation study, we use the Preqin data from Section 3.1 and form vintage-year portfolios (VYPs) to estimate simple linear SDF

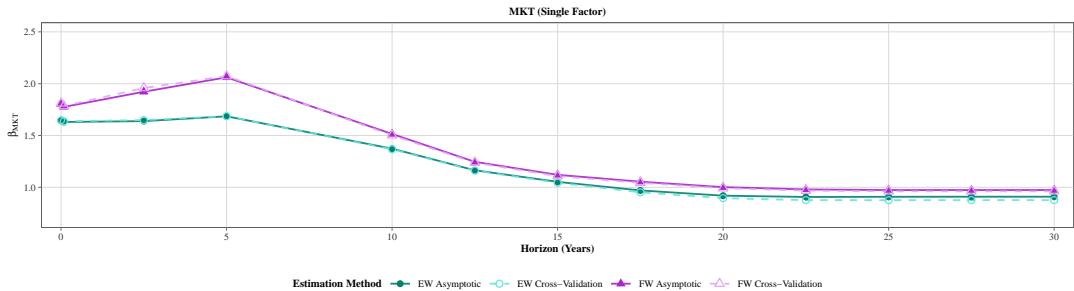


Figure 8: **Empirical (MKT factor):** Single-factor model for PE using VYP.

models with a Horizon of 15 years. For each model specification, we report both asymptotic standard errors (SE_A) and hv -block cross-validation standard errors (SE_{CV}). We focus the presentation on PE (i.e., mainly Buyout and Growth) as our primary fund type and report both equal-weighted (EW) and value-weighted (FW) estimates across a grid of compounding horizons. Throughout this subsection, t_A and t_{CV} denote the t -ratios based on asymptotic and cross-validation standard errors,¹¹.

3.4.1 Single market-factor models

Table 5 and Figure 8 report the single-factor MKT model for PE across horizons from 0.1 to 30 years. The market beta estimate broadly declines with the compounding horizon: under equal weighting (EW), $\hat{\beta}_{MKT}$ ranges from approximately 1.63–1.69 at short horizons (0–5 years) and decreases to approximately 0.91 at 25–30 years; under fund-size weighting (FW), the corresponding range is 1.78–2.06 down to 0.97. The elevated estimates at short horizons likely reflect the finite-sample bias documented in the simulation study, whereby the estimator overestimates market exposure when few discounting dates are used. Estimates stabilize beyond approximately 20 years, consistent with the bias-minimizing horizons identified in Section 3.3.

A striking feature of the results is the divergence between asymptotic and cross-validation inference. Asymptotic t -ratios (t_A) are borderline at short horizons (≈ 1.8) and drop well below 1.96 at longer horizons; partly driven by extremely large and erratic asymptotic standard errors (e.g., $SE_A > 1.7$ for horizons ≥ 15 years). In contrast, cross-validation t -ratios (t_{CV}) are uniformly large, ranging from 5.2 to 10.3 under EW and from 5.99 to 44.5 under FW, clearly indicating significance of the single-factor MKT model. Given the small number of cross-sectional units (vintage-year portfolios) available for estimation, the asymptotic approximation appears unreliable, and we regard the hv -block cross-validation results as more trustworthy. Therefore, drawing on both the simulation evidence and the cross-validation results, we focus on single MKT factor models for the remainder of the empirical analysis, even when their asymptotic t -ratios fall below conventional thresholds.

¹¹All R code and data are available in an online repository: https://github.com/quant-unit/Fundwise_SDF/tree/master/r_project.

Table 5: Empirical Estimates: PE — Single-Factor (MKT) Model

Horizon	Equal-Weighted (EW)					Value-Weighted (FW)				
	$\hat{\beta}$	SE_A	SE_{CV}	t_A	t_{CV}	$\hat{\beta}$	SE_A	SE_{CV}	t_A	t_{CV}
0 yr	1.642	0.900	0.160	1.82	10.31	1.801	1.011	0.256	1.78	7.08
5 yr	1.685	0.892	0.188	1.89	8.96	2.061	1.124	0.364	1.83	5.69
10 yr	1.373	0.864	0.172	1.59	7.97	1.515	0.959	0.259	1.58	5.80
15 yr	1.053	1.716	0.186	0.61	5.60	1.121	0.908	0.212	1.23	5.23
20 yr	0.920	1.856	0.172	0.50	5.22	1.002	1.744	0.214	0.57	4.63
30 yr	0.909	1.832	0.158	0.50	5.53	0.974	1.702	0.204	0.57	4.71

Notes: $\hat{\beta}$ = parameter estimate; SE_A = asymptotic standard error; SE_{CV} = cross-validation standard error; t_A/t_{CV} = corresponding t -ratios. Significance: * $p < 0.05$, ** $p < 0.01$, *** $p < 0.005$.

3.4.2 Two-factor models

We now examine two-factor models for PE that pair the MKT factor with either an Alpha intercept or one of the four non-market q^5 factors (Tables 6–10 and Figure 9).

MKT + Alpha (Table 6): Adding an α term increases the MKT loading relative to the single-factor model at most horizons: under EW, $\hat{\beta}_{MKT}$ ranges from 1.24 to 1.44; under FW, from 1.36 to 1.93. The Alpha coefficient is positive at short horizons ($\hat{\beta}_\alpha \approx 0.003$ –0.005 per month, corresponding to roughly 3.6–6% annualized) but turns negative at horizons beyond 10 years ($\hat{\beta}_\alpha \approx -0.002$ to -0.004 per month). This sign-reversal across horizons is consistent with the horizon-dependent bias pattern documented in the simulation study. While a few isolated asymptotic t -ratios are significant (e.g., EW at 5 years: $t_A = 10.8$ for MKT, $t_A = 3.49$ for Alpha), the asymptotic standard errors are highly erratic across horizons (SE_A ranging from 0.13 to over 26), precluding a reliable asymptotic inference for the two-factor model.

Non-market q^5 factors (Tables 7–10): A common pattern emerges across all four specifications (EG, IA, ME, ROE): the second-factor loading is positive at short horizons and declines (often changing sign) as the horizon increases, while the MKT loading absorbs a compensating adjustment. None of the second-factor coefficients are asymptotically significant at any horizon, and the asymptotic standard errors are again extremely large and unstable. Among the q^5 factors, ME stands out with particularly large positive loading estimates under FW at short horizons ($\hat{\beta}_{ME} \approx 1.0$ –1.6) that remain positive and CV-significant even at long horizons. In contrast, the EG, IA, and ROE loadings uniformly become negative beyond 10 years and remain insignificant under both weighting schemes.

Overall, the two-factor models for PE are fragile: estimates are highly horizon-dependent, asymptotic inference is unreliable, and no second factor consistently adds explanatory power beyond MKT. These findings corroborate the simulation-based recommendation to focus on single-factor MKT models given the limited cross-section of vintage-year portfolios currently available, avoiding over-interpretation of the non-market q^5 factor exposures.

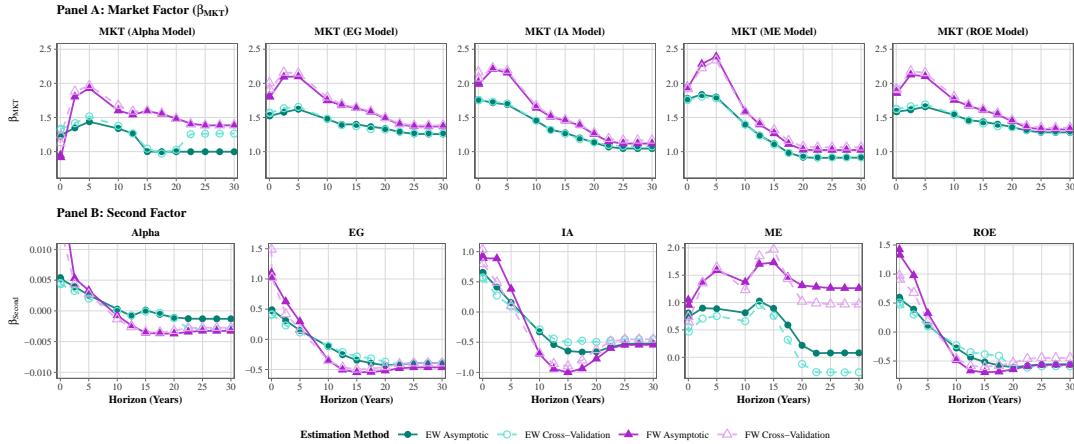


Figure 9: **Empirical (q^5 factors):** Two-factor models for PE using VYP.

Table 6: Empirical Estimates: PE — MKT + Alpha Model

Horizon	Equal-Weighted (EW)					Value-Weighted (FW)				
	$\hat{\beta}$	SE_A	SE_{CV}	t_A	t_{CV}	$\hat{\beta}$	SE_A	SE_{CV}	t_A	t_{CV}
<i>Panel A: Market Factor (β_{MKT})</i>										
0 yr	1.222	1.302	0.220	0.94	6.01	0.949	1.810	0.551	0.52	2.04
5 yr	1.438***	0.139	0.302	10.36	5.02	1.928	1.015	0.648	1.90	3.03
10 yr	1.338	12.857	0.374	0.10	3.68	1.601	24.401	0.613	0.07	2.74
15 yr	1.000	2.070	0.130	0.48	8.05	1.598	9.207	0.668	0.17	2.39
20 yr	0.999	2.131	0.164	0.47	6.27	1.488	6.691	0.748	0.22	1.98
30 yr	0.999	3.300	0.394	0.30	3.21	1.385	5.594	0.797	0.25	1.74
<i>Panel B: Alpha (β_{Alpha})</i>										
0 yr	0.005	0.007	0.002	0.73	2.40	0.020*	0.009	0.008	2.17	1.95
2.5 yr	0.004	0.003	0.002	1.19	1.39	0.005	0.010	0.004	0.56	1.05
5 yr	0.003***	0.001	0.002	3.51	1.01	0.003*	0.001	0.003	2.44	0.82
10 yr	0.000	0.038	0.003	0.01	-0.00	-0.001	0.158	0.004	-0.00	-0.34
15 yr	0.000	0.077	0.002	0.00	-0.07	-0.004	0.073	0.004	-0.05	-0.79
20 yr	-0.001	0.130	0.002	-0.01	-0.59	-0.004	0.062	0.005	-0.06	-0.65
30 yr	-0.001	0.211	0.003	-0.01	-0.88	-0.003	0.060	0.005	-0.05	-0.50

Notes: $\hat{\beta}$ = parameter estimate; SE_A = asymptotic standard error; SE_{CV} = cross-validation standard error; t_A/t_{CV} = corresponding t -ratios. Significance: * $p < 0.05$, ** $p < 0.01$, *** $p < 0.005$.

Table 7: Empirical Estimates: PE — MKT + EG Model

Horizon	Equal-Weighted (EW)					Value-Weighted (FW)				
	$\hat{\beta}$	SE_A	SE_{CV}	t_A	t_{CV}	$\hat{\beta}$	SE_A	SE_{CV}	t_A	t_{CV}
<i>Panel A: Market Factor (β_{MKT})</i>										
0 yr	1.528	1.818	0.250	0.84	6.28	1.817**	0.661	0.546	2.75	3.68
5 yr	1.621	17.286	0.282	0.09	5.87	2.101	4.304	0.448	0.49	4.80
10 yr	1.474	2.814	0.241	0.52	6.14	1.749	6.056	0.337	0.29	5.32
15 yr	1.400	1.279	0.216	1.09	6.34	1.637	2.492	0.289	0.66	5.70
20 yr	1.338	1.361	0.244	0.98	5.41	1.487	2.312	0.348	0.64	4.31
30 yr	1.257	1.439	0.277	0.87	4.59	1.368	1.832	0.373	0.75	3.72
<i>Panel B: EG (β_{EG})</i>										
0 yr	0.477	1.098	0.296	0.43	1.36	1.102**	0.407	1.838	2.71	0.81
5 yr	0.148	11.338	0.243	0.01	0.47	0.294	2.526	0.427	0.12	0.37
10 yr	-0.120	2.641	0.271	-0.05	-0.39	-0.348	4.373	0.368	-0.08	-0.95
15 yr	-0.345	4.565	0.318	-0.08	-0.89	-0.547	2.205	0.322	-0.25	-1.55
20 yr	-0.431	7.164	0.383	-0.06	-0.97	-0.520	2.175	0.366	-0.24	-1.24
30 yr	-0.392	8.473	0.289	-0.05	-1.36	-0.465	2.057	0.407	-0.23	-0.97

Notes: $\hat{\beta}$ = parameter estimate; SE_A = asymptotic standard error; SE_{CV} = cross-validation standard error; t_A/t_{CV} = corresponding t -ratios. Significance: * $p < 0.05$, ** $p < 0.01$, *** $p < 0.005$.

Table 8: Empirical Estimates: PE — MKT + IA Model

Horizon	Equal-Weighted (EW)					Value-Weighted (FW)				
	$\hat{\beta}$	SE_A	SE_{CV}	t_A	t_{CV}	$\hat{\beta}$	SE_A	SE_{CV}	t_A	t_{CV}
<i>Panel A: Market Factor (β_{MKT})</i>										
0 yr	1.763	1.222	0.280	1.44	6.27	2.015	1.754	0.567	1.15	3.82
5 yr	1.689	3.710	0.255	0.46	6.70	2.156	5.420	0.423	0.40	5.17
10 yr	1.451	1.126	0.206	1.29	7.06	1.640	5.551	0.282	0.30	5.92
15 yr	1.276	1.037	0.191	1.23	6.60	1.453	1.682	0.225	0.86	6.52
20 yr	1.132	1.001	0.240	1.13	4.74	1.255	1.776	0.299	0.71	4.26
30 yr	1.045	2.839	0.249	0.37	4.32	1.118	1.835	0.323	0.61	3.60
<i>Panel B: IA (β_{IA})</i>										
0 yr	0.656	2.240	0.537	0.29	1.07	0.932	3.059	1.448	0.30	0.72
5 yr	0.156	6.603	0.462	0.02	0.19	0.383	8.311	0.894	0.05	0.11
10 yr	-0.328	4.260	0.471	-0.08	-0.61	-0.696	10.654	0.619	-0.07	-1.10
15 yr	-0.646	3.722	0.553	-0.17	-0.91	-0.999	2.453	0.511	-0.41	-1.75
20 yr	-0.660	3.449	0.745	-0.19	-0.67	-0.774	2.986	0.713	-0.26	-0.88
30 yr	-0.525	4.910	0.727	-0.11	-0.62	-0.543	2.820	0.747	-0.19	-0.61

Notes: $\hat{\beta}$ = parameter estimate; SE_A = asymptotic standard error; SE_{CV} = cross-validation standard error; t_A/t_{CV} = corresponding t -ratios. Significance: * $p < 0.05$, ** $p < 0.01$, *** $p < 0.005$.

Table 9: Empirical Estimates: PE — MKT + ME Model

Horizon	Equal-Weighted (EW)					Value-Weighted (FW)				
	$\hat{\beta}$	SE_A	SE_{CV}	t_A	t_{CV}	$\hat{\beta}$	SE_A	SE_{CV}	t_A	t_{CV}
<i>Panel A: Market Factor (β_{MKT})</i>										
0 yr	1.781	1.143	0.217	1.56	8.13	1.955*	0.955	0.342	2.05	5.65
5 yr	1.800	1.142	0.218	1.58	8.16	2.392***	0.538	0.421	4.45	5.56
10 yr	1.401	1.008	0.191	1.39	7.31	1.593***	0.464	0.303	3.44	5.23
15 yr	1.118	0.939	0.187	1.19	5.90	1.267	1.579	0.255	0.80	5.16
20 yr	0.918	1.850	0.184	0.50	5.10	1.028	3.363	0.296	0.31	3.62
30 yr	0.912	1.830	0.180	0.50	5.09	1.024	3.133	0.288	0.33	3.69
<i>Panel B: ME (β_{ME})</i>										
0 yr	0.811	2.089	0.566	0.39	0.99	1.048	2.126	0.957	0.49	0.77
5 yr	0.888	2.234	0.296	0.40	2.55	1.587	3.536	0.529	0.45	3.11
10 yr	0.814	12.157	0.417	0.07	1.58	1.373	3.431	0.781	0.40	1.58
15 yr	0.893	2.958	0.346	0.30	2.19	1.732	6.762	0.412	0.26	4.78
20 yr	0.218	1.749	0.786	0.12	-0.15	1.313	7.376	1.023	0.18	1.00
30 yr	0.080	1.736	0.742	0.05	-0.37	1.263	9.526	0.989	0.13	0.98

Notes: $\hat{\beta}$ = parameter estimate; SE_A = asymptotic standard error; SE_{CV} = cross-validation standard error; t_A/t_{CV} = corresponding t -ratios. Significance: * $p < 0.05$, ** $p < 0.01$, *** $p < 0.005$.

Table 10: Empirical Estimates: PE — MKT + ROE Model

Horizon	Equal-Weighted (EW)					Value-Weighted (FW)				
	$\hat{\beta}$	SE_A	SE_{CV}	t_A	t_{CV}	$\hat{\beta}$	SE_A	SE_{CV}	t_A	t_{CV}
<i>Panel A: Market Factor (β_{MKT})</i>										
0 yr	1.588	2.540	0.248	0.63	6.57	1.877*	0.779	0.360	2.41	5.35
5 yr	1.657	18.700	0.261	0.09	6.49	2.103	12.191	0.421	0.17	5.12
10 yr	1.538	1.688	0.205	0.91	7.57	1.755	4.738	0.316	0.37	5.70
15 yr	1.437	1.213	0.179	1.18	7.90	1.605	2.205	0.273	0.73	5.94
20 yr	1.361	1.135	0.208	1.20	6.53	1.445	1.992	0.324	0.73	4.51
30 yr	1.282	1.134	0.239	1.13	5.41	1.325	1.843	0.338	0.72	4.00
<i>Panel B: ROE (β_{ROE})</i>										
0 yr	0.600	2.202	0.499	0.27	1.02	1.424*	0.699	0.933	2.04	1.05
5 yr	0.117	15.801	0.365	0.01	0.25	0.325	9.526	0.597	0.03	0.31
10 yr	-0.276	2.999	0.395	-0.09	-0.58	-0.489	4.149	0.516	-0.12	-0.89
15 yr	-0.516	2.828	0.536	-0.18	-0.72	-0.695	2.280	0.483	-0.30	-1.24
20 yr	-0.607	3.573	0.227	-0.17	-2.71	-0.645	2.145	0.551	-0.30	-0.96
30 yr	-0.558	3.742	0.243	-0.15	-2.43	-0.564	2.136	0.589	-0.26	-0.76

Notes: $\hat{\beta}$ = parameter estimate; SE_A = asymptotic standard error; SE_{CV} = cross-validation standard error; t_A/t_{CV} = corresponding t -ratios. Significance: * $p < 0.05$, ** $p < 0.01$, *** $p < 0.005$.

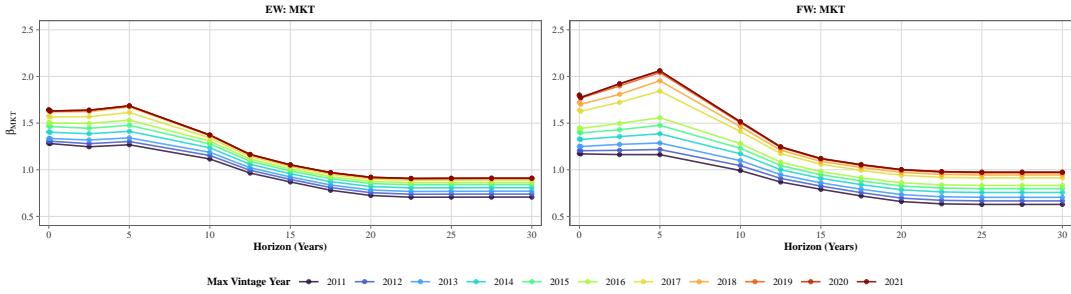


Figure 10: **Sensitivity analysis (MKT factor):** Single-factor model for PE using VYP with different vintage-year cutoffs (asymptotic estimates).

3.4.3 Focus on varying vintage cutoffs

Our standard vintage year cutoff in our empirical Preqin dataset is 2019. In this section, we analyze how altering the maximum vintage cutoff influences the parameter estimates. Concretely, we compare yearly cutoff-vintages from 2021 to 2011, making it 11 distinct cutoff dates; note that cutoffs beyond 2019 include only vintage years already present in the dataset but allow their most recent (and potentially still unrealized) cash flows to fully enter the estimation. We only consider asymptotic estimates in this section, i.e., we did not perform cross-validation results for the different vintage cutoffs.

Figure 10 shows the market beta estimates for the single-factor model across different vintage cutoffs. We observe a clear upward trend in the market beta as we include more recent vintages. For equal-weighted (EW) portfolios, $\hat{\beta}_{MKT}$ rises from approximately 0.87 with a 2011 cutoff to 1.05 with the full sample (2019/2021). Similarly, value-weighted (FW) estimates increase from 0.79 to 1.12. This could suggest that recent vintage years display higher market co-movements or that treating unrealized NAVs as distribution cash flows mechanically increases the β exposures. Interestingly, asymptotic standard errors tend to increase when including the most recent vintages (e.g., jumping from ≈ 0.5 to > 1.0 for FW after 2016), likely because these younger funds contribute fewer informative cash flows given the 15-year horizon.

Figure 11 displays the coefficients for the two-factor models. The second-factor estimates are largely robust to the vintage cutoff, maintaining their general sign and magnitude. For instance, the ME factor loading remains consistently positive (around 0.9 for EW and 1.7 for FW). In contrast, the Alpha coefficient in the MKT+Alpha model exhibits a downward drift, turning from insignificant positive values in the earlier cutoffs to significantly negative values (e.g., -0.18% per month) when including post-2016 vintages.

4 Conclusion

Theoretically, our Least-Mean-Distance estimator can be easily generalized to estimate SDF models for all kinds of non-traded cash flows. Practically, semiparametric

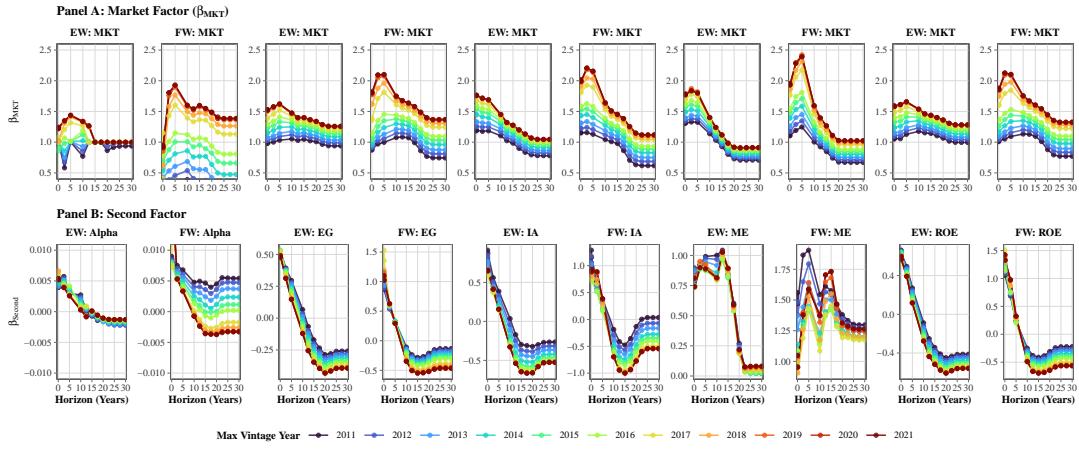


Figure 11: **Sensitivity analysis (q^5 factors):** Two-factor models for PE using VYP with different vintage-year cutoffs (asymptotic estimates).

estimators commonly exhibit problematic small sample behavior. Given the amount of currently available private equity fund data, our estimator's variance seems quite large, even for simple SDF model specifications. Specifically, our Monte Carlo simulation results prompt us to conclude that the closely related Driesssen et al. (2012) estimator may exhibit more bias and variance than originally assumed in their paper. Especially, the variance of α estimates seems to be too high to allow reliable abnormal performance conclusions. Fortunately, we show that the estimator's variance can be substantially lowered by averaging pricing errors over multiple dates within the fund lifetime, and that the small-sample bias is most effectively mitigated by increasing the cross-sectional sample size (n).

In the data-sparse private equity domain with only 20-40 cross-sectional units (i.e., vintage year portfolios) currently available for estimation, asymptotic inference seems not to be overly useful. Thus, we strongly advise always challenging asymptotic inference results by resampling or cross-validation techniques adapted to the dependence structure of overlapping fund cash flows. However, even these conclusions should be double-checked to avoid unreasonable instances, e.g., when hv -block cross-validation chooses dubious models with negative MKT factor estimates. Unfortunately, using individual funds instead of vintage year portfolios, which yields smaller asymptotic standard errors, constitutes no viable resolution as individual funds show considerably larger small-sample bias and variance in our Monte Carlo example (and substantially increased computational cost). Since, in our empirical analyses, basically all two-factor models' asymptotic standard errors appear statistically insignificant, we conjecture that naive versions of our SDF estimator shall be exclusively used for a single-MKT-factor model until considerably more vintage year information for private equity funds is available.

If someone wants to estimate more complex SDF models that incorporate additional factors, more structure is needed. These can be parametric assumptions for the data generating process (Ang et al., 2018) or to extract additional information from intermediate net asset values (Gredil et al., 2020; Brown et al., 2021).

A first “modern” approach to the same problem is applying statistical techniques incorporating sparse high-dimensional data (e.g., LASSO, ridge or model averaging) that regularize/shrink all coefficients other than the MKT factor. Secondly, given the high estimator variance revealed in the simulation study, statistical learning methods that create a strong learner by combining multiple weak learners seem also worth considering.

Another promising avenue for future research is extending the data generating process (DGP) to explicitly incorporate fee structures and carried interest. This would enable the estimation of SDFs that reflect gross-vs-net-of-fee performance and provide more precise evaluation models from both the General-Partner (GP) and Limited-Partner (LP) perspective. A more realistic DGP should naturally yield more accurate SDF parameter estimates and permit a sharper analytical characterization of the estimator’s bias and variance terms.

Additionally, treating intermediate net asset values (NAVs) as proxies for the unobservable true fund valuation could be explored to alleviate the timing risk premium bias associated with net future values. As a downside, such an approach would entail fundamentally adapting the estimator’s objective to place less emphasis on sparsely realized cash flows and rely more heavily on reported NAVs instead.

Finally, we point to the potentially most interesting topic for future research. Our simulation study indicates that the estimator’s bias and variance can be controlled by an appropriate choice for the set \mathcal{T} . This set averages the pricing error over multiple discounting dates. In simpler terms, an identification method that utilizes a future value concept instead of net present values obtains more favorable results in our case. The bias in our simulation study is minimal when the set of discounting dates corresponds approximately to the fund lifetime. A model that can determine the optimal value of \mathcal{T} for a given fund sample and macroeconomic environment would be highly appreciated.

References

- Andonov, A., Kräussl, R., and Rauh, J. (2021). Institutional investors and infrastructure investing. *The Review of Financial Studies*, 34(8):3880–3934.
- Andrews, D. W. (1991). Heteroskedasticity and autocorrelation consistent covariance matrix estimation. *Econometrica: Journal of the Econometric Society*, pages 817–858.
- Ang, A., Chen, B., Goetzmann, W. N., and Phalippou, L. (2018). Estimating private equity returns from limited partner cash flows. *Journal of Finance*, 73(4):1751–1783.
- Bertholon, H., Monfort, A., and Pegoraro, F. (2008). Econometric asset pricing modelling. *Journal of Financial Econometrics*, 6(4):407–458.
- Boyer, B. H., Nadauld, T. D., Vorkink, K. P., and Weisbach, M. S. (2023). Discount-rate risk in private equity: Evidence from secondary market transactions. *The Journal of Finance*, 78(2):835–885.

- Brown, G. W., Ghysels, E., and Gredil, O. (2021). Nowcasting net asset values: The case of private equity. working paper (as of 2021-03-19).
- Brown, G. W., Gredil, O. R., and Kaplan, S. N. (2019). Do private equity funds manipulate reported returns? *Journal of Financial Economics*, 132(2):267–297.
- Brown, G. W., Harris, R. S., Jenkinson, T., Kaplan, S. N., and Robinson, D. T. (2015). What do different commercial data sets tell us about private equity performance? Available at SSRN 2706556.
- Conley, T. G. (1999). Gmm estimation with cross-sectional dependence. *Journal of Econometrics*, 92(1):1–45.
- Cordell, L., Roberts, M. R., and Schwert, M. (2023). Clo performance. *The Journal of Finance*, 78(3):1235–1278.
- Driessen, J., Lin, T.-C., and Phalippou, L. (2012). A new method to estimate risk and return of nontraded assets from cash flows: the case of private equity. *Journal of Financial and Quantitative Analysis*, 47(3):511–535.
- Erel, I., Flanagan, T., and Weisbach, M. S. (2024). Risk-adjusting the returns to private debt funds. Technical report, National Bureau of Economic Research.
- Eu, C. N. L. (2017). *Numerical Analysis in Nonlinear Least Squares Methods and Applications*. PhD thesis, Curtin University.
- Fama, E. F. and French, K. R. (1993). Common risk factors in the returns on stocks and bonds. *Journal of financial economics*, 33(1):3–56.
- Farnsworth, H., Ferson, W., Jackson, D., and Todd, S. (2002). Performance evaluation with stochastic discount factors. *Journal of Business*, 75(3):473–503.
- Farrelly, K. and Stevenson, S. (2019). The risk and return of private equity real estate funds. *Global Finance Journal*, 42:100471.
- Franzoni, F., Nowak, E., and Phalippou, L. (2012). Private equity performance and liquidity risk. *The Journal of Finance*, 67(6):2341–2373.
- Giommetti, N. and Jørgensen, R. (2021). Risk adjustment of private equity cash flows. Available at SSRN 3980121.
- Gourieroux, C. and Monfort, A. (2007). Econometric specification of stochastic discount factor models. *Journal of Econometrics*, 136(2):509–530.
- Gredil, O., Sorensen, M., and Waller, W. (2020). Evaluating private equity performance using stochastic discount factors. working paper).
- Gupta, A. and Nieuwerburgh, S. V. (2021). Valuing private equity investments strip by strip. *Journal of Finance*.

- Han, C. and Phillips, P. C. (2006). Gmm with many moment conditions. *Econometrica*, 74(1):147–192.
- Hansen, B. E. and Racine, J. S. (2012). Jackknife model averaging. *Journal of Econometrics*, 167(1):38–46.
- Hansen, L. P. (1982). Large sample properties of generalized method of moments estimators. *Econometrica*, 50(4):1029–1054.
- Hansen, L. P. (2012). Proofs for large sample properties of generalized method of moments estimators. *Journal of Econometrics*, 170:325–330.
- Hansen, L. P. and Jagannathan, R. (1991). Implications of security market data for models of dynamic economies. *Journal of political economy*, 99(2):225–262.
- Hansen, L. P. and Richard, S. F. (1987). The role of conditioning information in deducing testable restrictions implied by dynamic asset pricing models. *Econometrica*, pages 587–613.
- Harris, R. S., Jenkinson, T., and Kaplan, S. N. (2014). Private equity performance: What do we know? *The Journal of Finance*, 69(5):1851–1882.
- Hou, K., Xue, C., and Zhang, L. (2015). Digesting anomalies: An investment approach. *Review of Financial Studies*, 28(3):650–705.
- Hou, K., Xue, C., and Zhang, L. (2020). Replicating anomalies. *The Review of Financial Studies*, 33(5):2019–2133.
- Hüther, N., Schmid, L., and Steri, R. (2022). Credit market equivalents and the valuation of private firms. Unpublished working paper.
- Jeffers, J., Lyu, T., and Posenau, K. (2024). The risk and return of impact investing funds. *Journal of Financial Economics*, 161:103928.
- Jenish, N. and Prucha, I. R. (2009). Central limit theorems and uniform laws of large numbers for arrays of random fields. *Journal of Econometrics*, 150(1):86–98.
- Jenish, N. and Prucha, I. R. (2012). On spatial processes and asymptotic inference under near-epoch dependence. *Journal of Econometrics*, 170(1):178–190.
- Jenkinson, T., Landsman, W. R., Rountree, B. R., and Soonawalla, K. (2020). Private equity net asset values and future cash flows. *The Accounting Review*, 95(1):191–210.
- Kim, M.-J., Kim, S.-S., and Lee, S.-H. (2015). Estimation of risk and return of venture capital investments in an emerging market: an iterative generalized method of moments approach. *Asia-Pacific Journal of Financial Studies*, 44(3):475–495.
- Kim, M. S. and Sun, Y. (2011). Spatial heteroskedasticity and autocorrelation consistent estimation of covariance matrix. *Journal of Econometrics*, 160(2):349–371.

- Korteweg, A. (2019). Risk adjustment in private equity returns. *Annual Review of Financial Economics*, 11:131–152.
- Korteweg, A. and Nagel, S. (2016). Risk-adjusting the returns to venture capital. *Journal of Finance*, 71(3):1437–1470.
- Korteweg, A. and Nagel, S. (2024). Risk-adjusted returns of private equity funds: a new approach.
- Korteweg, A. G. (2022). Risk and return in private equity. Handbook of the Economics of Corporate Finance, Vol 1: Private Equity and Entrepreneurial Finance, Forthcoming.
- Kroencke, T. A., Muehler, G., and Spiertsma, M. (2013). Return and risk of human capital contracts. *ZEW-Centre for European Economic Research Discussion Paper*, (13-108).
- Lee, J. M., Shilling, J. D., and Wurtzebach, C. (2019). A new method to estimate risk and return of commercial real estate assets from cash flows. *Journal of Real Estate Research*, 41(2):183–218.
- Leeb, H. and Pötscher, B. M. (2005). Model selection and inference: Facts and fiction. *Econometric Theory*, 21:21–59.
- Li, H., Xu, Y., and Zhang, X. (2016). Hedge funds performance evaluation under the stochastic discount factor framework. *Journal of Financial and Quantitative Analysis*, 51(1):231–257.
- McKinsey & Company (2024). McKinsey global private markets review 2024: Private markets: A slower era. Available online.
- Nagar, A. L. (1959). The bias and moment matrix of the general k-class estimators of the parameters in simultaneous equations. *Econometrica: Journal of the Econometric Society*, pages 575–595.
- Newey, W. K. (1993). Efficient estimation of models with conditional moment restrictions. In Engle, R. F. and McFadden, D. L., editors, *Handbook of Statistics*, volume 11, pages 419–454. Elsevier.
- Newey, W. K. and McFadden, D. (1994). Large sample estimation and hypothesis testing. *Handbook of econometrics*, 4:2111–2245.
- Newey, W. K. and Smith, R. J. (2004). Higher order properties of gmm and generalized empirical likelihood estimators. *Econometrica*, 72(1):219–255.
- Newey, W. K. and Windmeijer, F. (2009). Generalized method of moments with many weak moment conditions. *Econometrica*, 77(3):687–719.
- Nocedal, J. and Wright, S. (2006). *Numerical optimization*. Springer Science & Business Media.

- Pötscher, B. M. and Prucha, I. R. (1997). *Dynamic nonlinear econometric models: Asymptotic theory*. Springer Science & Business Media.
- Racine, J. (2000). Consistent cross-validatory model-selection for dependent data: hv-block cross-validation. *Journal of Econometrics*, 99:39–61.
- Riddiough, T. J. and Wiley, J. A. (2022). Private funds for ordinary people: Fees, flows, and performance. *Journal of Financial and Quantitative Analysis*, 57(8):3252–3280.
- Rilstone, P., Srivastava, V. K., and Ullah, A. (1996). The second-order bias and mean squared error of nonlinear estimators. *Journal of Econometrics*, 75(2):369–395.
- Tausch, C. and Pietz, M. (2024). Machine learning private equity returns. *The Journal of Finance and Data Science*, 10:100141.
- Zhao, J. (2020). Bottom-up portfolio construction of public return factors: Application to european private equity fund benchmarking. Master Thesis, Munich Business School.

A Bias–Variance Tradeoff of \mathcal{T} -Averaging

This appendix shows that \mathcal{T} -averaging introduces a small asymptotic bias via the Timing Risk Premium but achieves a net bias reduction in finite samples by lowering the variance of the pricing-error moments. We formalize this claim and reconcile it with the simulation results from Section 3.3.

A.1 Decomposition of $\mathbb{E}[\bar{\epsilon}_i(\theta_0)]$

Starting from the average pricing error (Equation 17) and the deal-level decomposition (Equation 3), we can write

$$\bar{\epsilon}_i(\theta_0) = \frac{1}{|\mathcal{T}_i|} \sum_{\tau \in \mathcal{T}_i} \epsilon_{\tau,i}(\theta_0) = \frac{1}{|\mathcal{T}_i|} \sum_{\tau \in \mathcal{T}_i} \sum_{j=1}^J \frac{\Psi_{t_{i,j}^{\text{Inv}}}}{\Psi_\tau} \delta_{i,j} \quad (25)$$

Taking the unconditional expectation yields

$$\mathbb{E}[\bar{\epsilon}_i(\theta_0)] = \frac{1}{|\mathcal{T}_i|} \sum_{\tau \in \mathcal{T}_i} \sum_{j=1}^J \mathbb{E}\left[\frac{\Psi_{t_{i,j}^{\text{Inv}}}}{\Psi_\tau} \delta_{i,j}\right] \quad (26)$$

For each (τ, j) pair, we decompose via the covariance identity:

$$\mathbb{E}\left[\frac{\Psi_{t_{i,j}^{\text{Inv}}}}{\Psi_\tau} \delta_{i,j}\right] = \underbrace{\text{Cov}\left(\frac{\Psi_{t_{i,j}^{\text{Inv}}}}{\Psi_\tau}, \delta_{i,j}\right)}_{\text{TRP}_{\tau,j}} + \underbrace{\mathbb{E}\left[\frac{\Psi_{t_{i,j}^{\text{Inv}}}}{\Psi_\tau}\right]}_{>0} \cdot \underbrace{\mathbb{E}[\delta_{i,j}]}_{=0} \quad (27)$$

Hence the asymptotic bias of the average pricing error is entirely determined by the averaged Timing Risk Premium:

$$\mathbb{E}[\bar{\epsilon}_i(\theta_0)] = \frac{1}{|\mathcal{T}_i|} \sum_{\tau \in \mathcal{T}_i} \sum_{j=1}^J \text{Cov}\left(\frac{\Psi_{t_{i,j}^{\text{Inv}}}}{\Psi_\tau}, \delta_{i,j}\right) \quad (28)$$

The crucial observation is that this covariance is exactly zero when $\tau \leq t_{i,j}^{\text{Inv}}$ (Case A from Section 2.1). When $\tau \leq t_{i,j}^{\text{Inv}}$, iterated expectations through $\mathcal{F}_{t_{i,j}^{\text{Inv}}}$ give

$$\text{Cov}\left(\frac{\Psi_{t_{i,j}^{\text{Inv}}}}{\Psi_\tau}, \delta_{i,j}\right) = \mathbb{E}\left[\frac{\Psi_{t_{i,j}^{\text{Inv}}}}{\Psi_\tau} \underbrace{\mathbb{E}\left[\delta_{i,j} \mid \mathcal{F}_{t_{i,j}^{\text{Inv}}}\right]}_{=0 \text{ by Eq. 2}}\right] = 0 \quad (29)$$

Therefore the bias decomposes as

$$\mathbb{E}[\bar{\epsilon}_i(\theta_0)] = \frac{1}{|\mathcal{T}_i|} \left[\underbrace{\sum_{\substack{\tau \in \mathcal{T}_i, j \\ \tau \leq t_{i,j}^{\text{Inv}}}} 0}_{\text{NPV terms}} + \underbrace{\sum_{\substack{\tau \in \mathcal{T}_i, j \\ \tau > t_{i,j}^{\text{Inv}}}} \text{Cov}\left(\frac{\Psi_{t_{i,j}^{\text{Inv}}}}{\Psi_\tau}, \delta_{i,j}\right)}_{\text{NFV terms (Timing Risk Premium)}} \right] \quad (30)$$

Remark (misspecification). Under model misspecification, $\mathbb{E}[\delta_{i,j}] \neq 0$ and the second term in Equation 27 no longer vanishes. This introduces an additional bias component proportional to $\mathbb{E}[\Psi_{t_{i,j}^{\text{Inv}}} / \Psi_\tau] \cdot \mathbb{E}[\delta_{i,j}]$ that grows with $|\mathcal{T}_i|$.

A.2 Reconciliation with the simulation study

If this asymptotic bias were the only effect, the NPV-only approach of Driessen et al. (2012) with $\mathcal{T}_i = \{\min_j t_{i,j}^{\text{Inv}}\}$ would exhibit the smallest bias, and adding NFV terms would always increase it. However, the simulation study in Section 3.3 shows the opposite: averaging over larger \mathcal{T}_i reduces total bias. The resolution involves two competing effects.

Effect 1: Finite-sample bias from nonlinear estimation For the nonlinear least-squares objective $Q_n(\theta) = -\frac{1}{n} \sum_{i=1}^n \bar{\epsilon}_i(\theta)^2$, the finite-sample bias of $\hat{\theta}$ is approximately (standard Nagar (1959)-type expansion for M-estimators):

$$\mathbb{E}[\hat{\theta} - \theta_0] \approx -\frac{1}{n} H^{-1} \cdot b(\text{Var}[\bar{\epsilon}_i], \text{curvature}) \quad (31)$$

where $H = \nabla^2 Q_n(\theta_0)$ denotes the Hessian of the objective function. The term curvature refers to the second-order derivatives $\nabla_\theta^2 \bar{\epsilon}_i(\theta_0)$ of the moment conditions with respect to the parameter vector. These derivatives measure how nonlinear the estimating equations are in θ : if the pricing-error moments were perfectly linear in θ , the curvature would vanish and no finite-sample bias would arise. The correction

term $b(\cdot)$ is the Nagar-type bias correction from Rilstone et al. (1996) and Newey and Smith (2004); it captures the interaction between the variance of the moment conditions $\text{Var}[\bar{\epsilon}_i(\theta_0)]$ and the curvature of the objective. Specifically, b is proportional to $\text{Var}[\bar{\epsilon}_i(\theta_0)]$: the noisier the moments, the more the nonlinear optimizer is “pulled away” from the true θ_0 in finite samples, because the concavity/convexity of Q_n causes asymmetric responses to positive and negative moment deviations. When each fund provides only one noisy moment (NPV-only), this variance is high, resulting in a substantial finite-sample bias.

Effect 2: Variance reduction through \mathcal{T} -averaging When averaging over $|\mathcal{T}_i|$ discounting dates, the variance of the average pricing error is

$$\text{Var}[\bar{\epsilon}_i] = \frac{1}{|\mathcal{T}_i|^2} \sum_{\tau, \tau' \in \mathcal{T}_i} \text{Cov}(\epsilon_{\tau,i}, \epsilon_{\tau',i}) \quad (32)$$

The pricing errors $\epsilon_{\tau,i}$ at different τ are positively correlated (they share the same deal cash flows and factor realizations), but they are not perfectly correlated because different τ ’s weight the SDF path differently, so that exploiting these additional conditional moment restrictions improves efficiency (Newey, 1993). Hence

$$\text{Var}[\bar{\epsilon}_i] < \text{Var}[\epsilon_{\tau_0,i}] \quad \text{for } |\mathcal{T}_i| > 1$$

and this variance reduction directly decreases the finite-sample bias from the non-linear optimization (cf. Equation 31).

The net effect The total bias in small samples decomposes approximately as

$$\underbrace{\text{Finite-sample bias}}_{\propto \text{Var}[\bar{\epsilon}_i]/n} + \underbrace{\text{Asymptotic TRP bias}}_{\text{Eq. 30}}$$

As $|\mathcal{T}_i|$ increases from one:

- The finite-sample bias *decreases* because $\text{Var}[\bar{\epsilon}_i]$ drops (cf. Equation 31).
- The TRP bias *increases* as more NFV terms enter, but is tempered by two factors:
 - only (τ, j) pairs with $\tau > t_{i,j}^{\text{Inv}}$ contribute non-zero TRP terms,
 - the $1/|\mathcal{T}_i|$ denominator dilutes the accumulating terms.

The simulation study finds the sweet spot where these two effects balance: at $|\mathcal{T}_i|$ corresponding to approximately the fund lifetime (180 months), the variance-reduction benefit dominates the TRP cost.

A.3 Interpretation

The \mathcal{T} -averaging mechanism can be understood through three complementary perspectives from the econometrics and asset pricing literature.

Variance reduction via moment combination. In the GMM framework of Hansen (1982), efficiency gains arise from combining multiple moment conditions. Each discounting date $\tau \in \mathcal{T}_i$ contributes a separate pricing-error moment $\epsilon_{\tau,i}(\theta)$, and averaging over \mathcal{T}_i is formally equivalent to combining $|\mathcal{T}_i|$ such moments into a single summary condition. Since these moments share the same underlying deal cash flows and factor paths but weight the SDF trajectory differently, they are positively yet imperfectly correlated. The resulting reduction in $\text{Var}[\bar{\epsilon}_i]$ is the classical Hansen (1982) effect of exploiting additional moment information, which in turn decreases the finite-sample bias of the nonlinear estimator (cf. Equation 31).

The TRP as a bounded risk-premium term. The Timing Risk Premium in Equation 28 is a covariance between a ratio of stochastic discount factors and a deal-level pricing error. Within the Hansen and Jagannathan (1991) framework, the magnitude of such covariances is bounded by the product of the SDF’s variability and the payoff’s standard deviation. Under correct model specification, the TRP terms are therefore non-zero (i.e., they reflect genuine systematic risk) but their size is disciplined by the Hansen–Jagannathan bound. Consequently, the TRP bias introduced by NFV terms ($\tau > t_{i,j}^{\text{Inv}}$) grows only moderately with $|\mathcal{T}_i|$ and is further diluted by the $1/|\mathcal{T}_i|$ denominator.

Analogy to model averaging. The structure of \mathcal{T} -averaging closely mirrors the jackknife model averaging framework of Hansen and Racine (2012), where combining multiple (possibly misspecified) estimators achieves a bias–variance tradeoff: averaging reduces variance at the expense of a small bias from the inclusion of imperfect components. In our setting, the “models” being averaged are the pricing-error moments at different discounting dates τ . Each individual moment is unbiased for the NPV terms but may carry TRP bias for the NFV terms; analogous to the weak misspecification in the Hansen and Racine (2012) setup. Their key insight applies directly: the variance reduction from averaging dominates the bias cost as long as the individual components are not too severely biased, which is precisely what the Hansen and Jagannathan (1991) bound guarantees in our context.

Summary. The \mathcal{T} -averaging mechanism does not make the TRP wash out in the sense that each individual TRP term goes to zero. Rather, it achieves a bias–variance tradeoff where the TRP bias introduced by NFV terms is more than compensated by the variance reduction that cures the finite-sample bias from the nonlinear estimation. The fact that the optimal $|\mathcal{T}_i|$ depends on the SDF model specification is consistent with this explanation: a more complex or misspecified model not only increases the TRP covariance terms but also activates the additional bias channel identified in the Remark on misspecification above, shifting the optimal tradeoff point. Finally, the favorable tradeoff observed in our simulation study may partly reflect the conditioning on a single realized factor path, which attenuates the TRP covariance terms relative to their unconditional magnitude. Across different macro regimes, the TRP cost could be larger and the optimal $|\mathcal{T}_i|$ smaller.

B Comparison to Driessen et al. (2012)

B.1 Comparison to simulation results

(Driessen et al., 2012, Table 1) report Monte Carlo results for $n = 50$ funds per vintage-year portfolio with an annualized idiosyncratic volatility of 108% ($\sigma_{\text{monthly}} = 0.3$). In our base case (S1), we use only $n = 20$ funds per vintage, which is more realistic for older vintages and smaller private capital asset classes beyond Buyout and Venture Capital. Scenarios S2–S9 in Table 11 systematically vary the simulation design to identify which features drive the remaining discrepancy between our results and the Driessen et al. (2012) benchmark.

Several structural differences between the two simulation designs should be noted. First, Driessen et al. (2012) use a shifted lognormal distribution for the idiosyncratic error term (cf. their online appendix, Equation 9), which ensures that the error addend has a finite lower bound (approximately -1) and avoids the non-linear distortions that arise when normally distributed errors occasionally produce extreme negative total returns. Second, they employ a “full Monte Carlo” approach and simulate the total market return R_{MKT} from a shifted lognormal distribution calibrated to S&P 500 moments (1980–2003), rather than conditioning on the historical path. Third, their deal investment timing follows a deterministic schedule (three deals per year, evenly spaced across the five-year investment period), rather than random uniform draws.

Interpretation of results. Under our original design with normal errors and $\sigma = 30\%$ (S3), the market beta is severely attenuated ($\hat{\beta} = 0.62$) and the spurious alpha (0.68%/month) is an order of magnitude larger than the benchmark (0.05%). Switching to the shifted lognormal error distribution (S4) modestly improves the beta estimate to 0.71 and reduces the alpha bias to 0.57%. Adding simulated market returns (S5) yields no further improvement ($\hat{\beta} = 0.66$), and deterministic investment timing (S7) produces a similarly attenuated beta of 0.70 without substantially reducing alpha variance.

The most effective modification is shortening the maximum deal holding period from 10 years to 5 years (S8), which significantly improves the beta estimate ($\hat{\beta} = 0.91$) and reduces the alpha bias to 0.30% without requiring lower idiosyncratic volatility. Combining this holding-period restriction with deterministic investment timing (S9) yields a similar beta estimate ($\hat{\beta} = 0.88$) and alpha bias (0.32%), with the deterministic timing contributing only a modest additional effect. By limiting the compounding duration, both scenarios restrict the geometric growth of idiosyncratic errors, reducing the frequency of degenerate alpha estimates and partially restoring the accuracy of the market beta. However, even the most favorable holding-period scenario (S8) still exhibits substantially more beta attenuation ($\hat{\beta} = 0.91$ vs. 0.98) and alpha bias (0.30% vs. 0.05%) than the Driessen et al. (2012) benchmark. Furthermore, tightening the alpha upper bound $\bar{\alpha}$ dramatically affects results: at $\bar{\alpha} = 0.5\%$ (S6), 32% of draws hit the bound, and beta is close to the benchmark ($\hat{\beta} = 0.87$) — but at the cost of artificially constraining the alpha search space. This sensitivity reveals a fundamental tension in the two-factor model: the alpha term absorbs residual nonlinearity that the linear SDF cannot capture, and bounding it arbitrarily distorts the beta estimate.

Overall, the large remaining discrepancy between our simulation results and those of Driessen et al. (2012) — even after adopting their shifted lognormal errors and simulated market returns — suggests that the favorable small-sample properties reported in Driessen et al. (2012) may be overly optimistic when the alpha upper bound is set more liberally ($\bar{\alpha} = 2\%$ in our default configuration vs. $\bar{\alpha} = 1\%$ in Driessen et al. (2012)). Across all scenarios with $\sigma = 30\%$, the beta estimates range from 0.62 to 0.91, well below the Driessen et al. (2012) benchmark of 0.98, confirming the estimator’s sensitivity to the interplay between idiosyncratic noise, compounding horizon, and alpha bounds.

Table 11: Comparison to simulation (Driessen et al., 2012, Table 1)

Scenario	Market Beta (β)						Alpha (α , monthly %)				
	True	Mean	SD	25%	75%	True	Mean	SD	25%	75%	% Bound
(S1) $n = 20$, $\sigma = 20\%$, Normal	1.00	0.84	0.40	0.61	1.05	0.00	0.22	0.28	0.04	0.32	0%
(S2) $n = 50$, $\sigma = 20\%$, Normal	1.00	0.89	0.29	0.74	1.03	0.00	0.13	0.16	0.03	0.19	0%
(S3) $n = 50$, $\sigma = 30\%$, Normal	1.00	0.62	0.72	0.12	1.03	0.00	0.68	0.70	0.13	1.08	0%
(S4) $n = 50$, $\sigma = 30\%$, Shifted LN	1.00	0.71	0.72	0.23	1.12	0.00	0.57	0.67	0.07	0.91	0%
(S5) $n = 50$, $\sigma = 30\%$, Shifted LN, Sim. MKT	1.00	0.66	0.79	0.11	1.06	0.00	0.71	0.79	0.03	1.34	0%
(S6) $n = 50$, $\sigma = 30\%$, Shifted LN, Sim. MKT, $\bar{\alpha} = 0.5\%$	1.00	0.87	0.69	0.41	1.24	0.00	0.48	0.50	0.06	1.00	32%
(S7) $n = 50$, $\sigma = 30\%$, Shifted LN, Sim. MKT, Det. Timing	1.00	0.70	0.79	0.19	1.12	0.00	0.65	0.80	-0.01	1.20	0%
(S8) $n = 50$, $\sigma = 30\%$, Shifted LN, Sim. MKT, Max. Hold 5	1.00	0.91	0.54	0.64	1.19	0.00	0.30	0.52	-0.03	0.47	0%
(S9) $n = 50$, $\sigma = 30\%$, Shifted LN, Sim. MKT, Max. Hold 5, Det. Timing	1.00	0.88	0.54	0.61	1.16	0.00	0.32	0.51	-0.01	0.53	0%
<i>Driessen et al. 2012, Table 1</i>											
Benchmark Estimates ($n = 50$, $\sigma_{\text{monthly}} = 30\%$)	1.00	0.98	0.32	0.80	1.15	0.00	0.05	0.2	-0.06	0.15	n/a

Notes: This table reports the small-sample properties of the generalized estimator for Horizon 0 (i.e., treating funds as held to maturity), mimicking the structure of Table 1 in Driessen et al. (2012). We simulate 1,000 Monte Carlo draws. All scenarios use 20 vintage years (1986–2005) with 15 deals per fund; unless stated otherwise, $n = 20$ funds per vintage, $\sigma = 0.2$ monthly idiosyncratic volatility, and a normal error distribution. The simulation dimensions vary as follows: n = number of funds per vintage-year portfolio; σ = monthly idiosyncratic standard deviation; *Normal* = Gaussian error term (default); *Shifted LN* = shifted lognormal error term (cf. Driessen et al. (2012)); *Sim. MKT* = total market return drawn from a shifted lognormal (S&P 500 1980–2003 moments) with a constant risk-free rate of 4% annualized, instead of conditioning on the historical path; *Det. Timing* = deterministic (evenly spaced) deal investment dates instead of random; *Max. Hold 5* = maximum holding period of a deal limited to 5 years (default 10 years); $\bar{\alpha}$ = upper bound on the alpha parameter in the optimization (default $\bar{\alpha} = 2\%/\text{month}$). The *% Bound* column reports the fraction of draws where the alpha estimate reached the upper optimization bound $\bar{\alpha}$. The final row presents the Benchmark Estimates obtained directly from Table 1 in Driessen et al. (2012) for comparative purposes. Alpha estimates are monthly and presented in percentages.

B.2 Comparison to empirical results

We now compare our empirical results directly with those published by Driessen et al. (2012). A key distinction is the underlying private market dataset: we use the Prequin cash flow database, which is also employed by Korteweg and Nagel (2016) and Ang et al. (2018), whereas Driessen et al. (2012) rely on the Thomson Venture Economics (TVE) fund-level dataset, which suffers from well-documented data quality issues such as forward-filling of stale KPIs (Harris et al., 2014). (Driessen et al., 2012, Table 5) report results for Buyout (BO) and Venture Capital (VC) funds using both an S&P 500 single-factor model and the Fama and French (1993) three-factor model. In their single-factor specification, VC exhibits a market beta

of 2.73 and a highly negative annualized alpha of -12.3% , while BO yields a market beta of 1.31 and an annualized alpha of -4.8% . These strongly negative alpha estimates can be at least partially attributed to the aforementioned data errors in TVE; indeed, both alpha coefficients are statistically insignificant once bootstrap standard errors are taken into account. In their three-factor specification, the factor loadings are broadly plausible: BO exhibits $\hat{\beta}_{MKT} = 0.72$, $\hat{\beta}_{SMB} = 0.50$, and $\hat{\beta}_{HML} = 0.19$, while VC shows $\hat{\beta}_{MKT} = 2.20$, $\hat{\beta}_{SMB} = 1.83$, and $\hat{\beta}_{HML} = -0.48$. However, Driessen et al. (2012) note that the three-factor model does not significantly improve upon the single-factor specification and that the SMB and HML loadings are largely insignificant.

Table 12 presents our own three-factor (MKT, SMB, HML) and four-factor (α , MKT, SMB, HML) Fama-French model results for BO, VC, PE, and RE at Horizon 0. The three-factor results in Panel A (EW) and Panel B (FW) confirm the qualitative pattern reported by Driessen et al. (2012): market betas are consistently above one, ranging from 1.23 (RE, EW) to 2.29 (VC, EW), broadly in line with the Driessen et al. (2012) estimates. The SMB loading is positive for BO, PE, and RE but negative for VC under both weighting schemes. Notably, this sign for VC *reverses* the pattern reported by Driessen et al. (2012), who find a strongly positive VC SMB loading ($\hat{\beta}_{SMB} = 1.83$). The discrepancy likely reflects differences in the underlying datasets (Preqin vs. TVE) and sample periods, and underscores the fragility of multi-factor estimates in private equity. However, the asymptotic standard errors for all three-factor loadings are extremely large — often exceeding the point estimates by a factor of five or more (e.g., $SE_{SMB} = 11.8$ for BO under EW, $SE_{SMB} = 58.1$ for BO under FW) — rendering every coefficient statistically insignificant. This imprecision is dramatically worse than the bootstrap standard errors reported in (Driessen et al., 2012, Table 5), where at least the MKT loading achieves marginal significance for some specifications.

The four-factor model results, which additionally include a monthly α term constrained to an upper bound of 2% per month, are even more problematic. The alpha estimates for VC and PE under FW hit the upper bound ($\hat{\alpha} = 0.020$), signaling that the optimizer is unable to find an interior solution — a symptom of the “exploding alpha” problem documented in our simulation study. Because the alpha term absorbs the compounded idiosyncratic noise, the remaining factor loadings become unstable: the MKT coefficient for BO drops to 0.020 under FW (essentially zero), yet SMB and HML loadings swing wildly (e.g., $\hat{\beta}_{HML} = -1.64$ for BO under FW). Strikingly, the asymptotic standard errors for MKT, SMB, and HML in the four-factor model round to (0.000), a numerical artifact of near-degenerate Hessian matrices caused by the alpha-induced collinearity. These findings reinforce our simulation-based conclusion from Section 3.3: with the limited cross-section of vintage-year portfolios currently available, multi-factor models are severely over-parameterized, and the single-factor MKT specification remains the only reliably estimable model.

Table 12: Estimation Results: 3-Factor vs 4-Factor Models

Model	Factor	Asset Class			
		BO	VC	PE	RE
<i>Panel A: Equal-Weighted (EW)</i>					
3-Factor	MKT	1.794 (2.654)	2.285 (1.317)	1.805 (1.403)	1.230 (2.136)
	SMB	1.431 (11.822)	-0.430 (5.504)	0.974 (8.947)	1.527 (3.158)
	HML	0.222 (5.032)	0.185 (1.005)	0.175 (4.653)	0.203 (7.995)
	Alpha	0.007 (0.001)	0.020 (0.001)	0.005 (0.003)	0.009 (0.001)
	MKT	1.015 (0.000)	1.052 (0.000)	1.260 (0.000)	0.308 (0.000)
	SMB	0.006 (0.000)	0.230 (0.000)	0.048 (0.000)	-0.005 (0.000)
4-Factor	HML	0.003 (0.000)	0.279 (0.000)	-0.059 (0.000)	-0.140 (0.000)
	<i>Panel B: Fund-Weighted (FW)</i>				
	MKT	1.809 (13.798)	2.150 (0.473)	1.835 (11.657)	1.430 (3.732)
	SMB	1.872 (58.064)	-0.703 (2.269)	1.926 (51.756)	2.786 (8.819)
	HML	-0.016 (11.580)	-0.097 (1.250)	-0.072 (22.261)	0.355 (2.767)
	Alpha	0.017 (0.010)	0.020 (0.001)	0.020 (0.002)	0.013 (0.004)
4-Factor	MKT	0.020 (0.000)	1.234 (0.000)	0.991 (0.000)	0.152 (0.000)
	SMB	1.415 (0.000)	0.484 (0.000)	-1.388 (0.000)	2.240 (0.000)
	HML	-1.648 (0.000)	0.696 (0.000)	-0.553 (0.000)	-0.869 (0.000)

Note: This table presents the estimation results of the Fama-French 3-factor (MKT, SMB, HML) and 4-factor (including α) models for different asset classes. Cash flows are discounted to horizon 0. Standard errors are reported in parentheses below the estimates. The models are estimated using the Least-Mean-Distance method. Panel A uses equal-weighted (EW) cash flows, while Panel B uses fund-size-weighted (FW) cash flows. The upper bound for α was constrained to 0.02.